



**aws**   
certified

**Solutions  
Architect**

Associate



# EC2 - Elastic Compute Cloud

## Types of IP

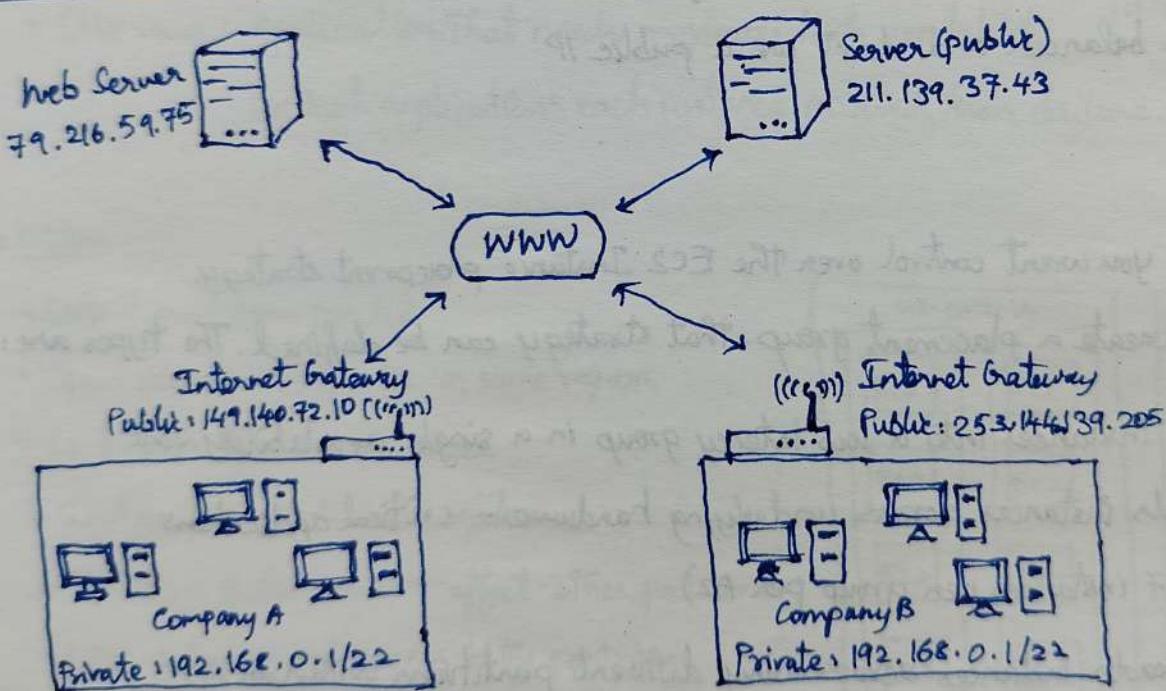
### a) IPv4

- Most commonly used format.
- Eg: 1.160.10.240
- Allows 3.7 billion different addresses in the public space.
- Range: [0-255].[0-255].[0-255].[0-255]

### b) IPv6

- Newest and solves problems for the Internet of Things (IoT)
- Eg: 3ffe:1900:4545:3:200:ff8fff:fe21:67ef

## Private vs Public IP



## Public IP

- Can be identified on the internet
- Must be unique across whole web
- Two machines cannot have same IP
- Can be geo-located easily

## Private IP

- Can be identified on private network only
- Must be unique across private network
- 2 different private networks (2 companies) can have same IPs
- Machines connect to www using Internet gateway
- Only specified range can be used as private IP

## Elastic IPs

- When you stop and then start an EC2 instance, it can change its public IP
- If you need to have a fixed public IP for your instance, you need an Elastic IP
- It is a public IPv4 IP, you own as long as you don't delete it.
- You can attach to one instance at a time only
- You can mask the failure of an instance or software by rapidly remapping the address to another instance in your account.
- You can have 5 Elastic IP in your account (you can ask AWS to increase)

## Avoid using Elastic IP:

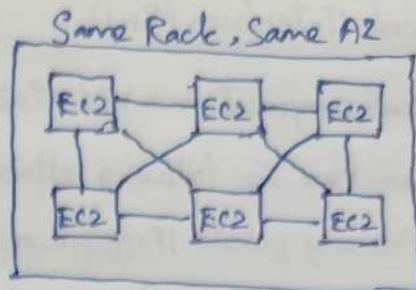
- They often reflect poor architectural decisions.
- Instead, use a random public IP and register a DNS name to it.
- Use a load balancer and don't use a public IP

## Placement Groups

- Sometimes, you want control over the EC2 Instance placement strategy.
- When you create a placement group, that strategy can be defined. The types are:
  - a) Cluster → clusters instances into a low latency group in a single availability zone.
  - b) Spread → spreads instances across underlying hardware - critical applications  
(max 7 instances per group per AZ)
  - c) Partition → spreads instances across many different partitions within an AZ  
Scales to 100s of EC2 instances per group (Hadoop, Cassandra, Kafka)

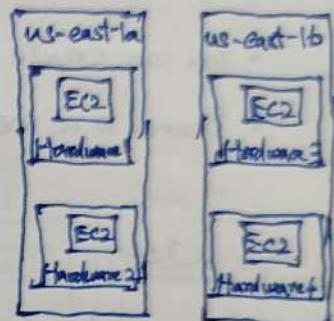
## a) Cluster

- Pros: Great network (10 Gbps bandwidth)
- Cons: If rack fails, all instances fail
- Use case:
  - Big data job that needs to complete fast
  - Application that needs extremely low latency and high network throughput



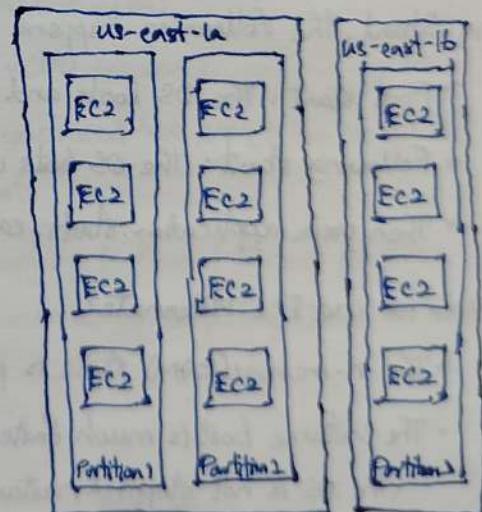
## b) Spread

- Pros: Span across multiple AZ
- Reduced risk of simultaneous failure
- EC2 instances are on different physical hardware
- Cons: Limited to 7 instances per AZ per placement group
- Use case: Application that needs maximize high availability
  - Critical applications each instance isolated from failure.



## c) Partition

- Up to 7 partitions per AZ
- Span across multiple AZ in same region
- Up to 100s of EC2 instances
- Instances do not share racks
- Partition failure won't affect other partitions
- EC2 instances get access to the partition information as metadata
- Use cases: HDFS, HBase, Cassandra, Kafka



## Elastic Network Interfaces (ENI)

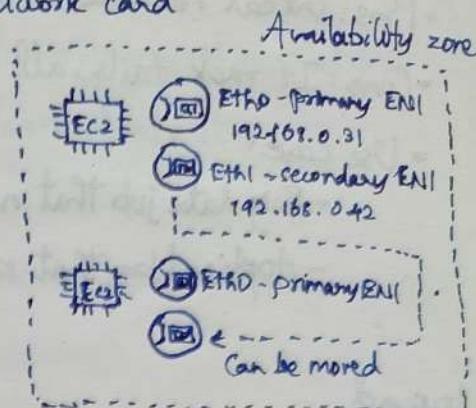
\* Logical component in a VPC that represents a virtual network card

\* It can have the following attributes:

- Primary private IPv4, one or more secondary IPv4
- One Elastic IP (IPv4) per private IPv4
- One Public IPv4
- One or more security groups
- A MAC address

\* You can create ENI independently and attach them on the fly (more them) on EC2 instances (to multiple instances)

\* Bound to a specific availability zone (AZ)



## EC2 Hibernate

\* We can stop, terminate instances

\* Stop : The data on disk (EBS) is kept intact in the next start

\* Terminate : any EBS volumes (root) also set-up to be destroyed is lost

\* On Start, the following happens

\* First start : The OS boots and the EC2 User Data script is run

\* Following start : The OS boots up

\* Then your application starts - caches get warmed up and that can take time

\* When we use EC2 Hibernate :

\* The in-memory (RAM) state is preserved

\* The instance boot is much faster

(the OS is not stopped/restarted)

\* RAM state is written to a file in root EBS volume

\* The root EBS volume must be encrypted

\* Use Cases:

- \* long-running processing
- \* saving the RAM state
- \* services that take time to initialize

### Note:

\* Instance RAM size < 150 GiB

\* Available for on-demand & Reserved instances

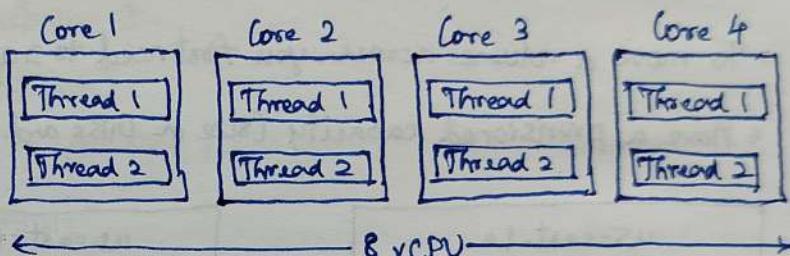
\* Instance cannot be hibernated > 60 days

## EC2 Nitro

- Underlying platform for the next generation of EC2 instances.
- New virtualization technology.
- Allows for better performance
  - Better networking options (enhanced networking, HPC, IPv6)
  - Higher Speed EBS (Nitro is necessary for 64000 EBS IOPS - max 32000 on non-Nitro)
- Better underlying security.

## EC2 vCPU

- Multiple threads can run on one CPU (multithreading)
- Each thread is represented as a virtual CPU (vCPU)
- Eg: m5.2x large
  - 4 CPU
  - 2 threads per CPU
  - 8 vCPU in total



## Optimizing CPU

- EC2 instances come with a combination of RAM and vCPU
- But in some cases, you may want to change the vCPU options:
  - # of CPU cores : you can decrease it (if you need high RAM & low No. of CPU)
    - to decrease licensing costs
  - # of threads per core : disable multithreading to have 1 thread per CPU
    - helpful for high performance computing (HPC) workloads
- Only specified during instance launch.

## EC2 Capacity Reservations

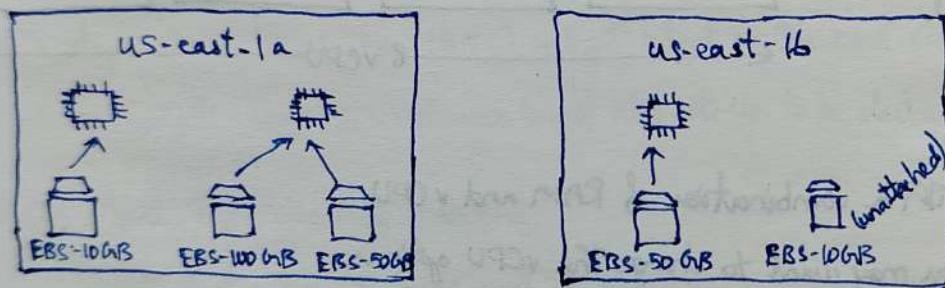
- Ensure you have EC2 capacity when needed
- Manual or planned end-date for reservation
- No need for 1 or 3-year commitment
- Capacity access is immediate
- You get billed as soon as it starts
- Combine with Reserved Instances and Savings plans to do cost saving.

### Specify:

- Availability zone to reserve capacity.
- No. of instances to reserve capacity.
- Instance attributes, instance type, tenancy and platform / OS

## EBS Volume

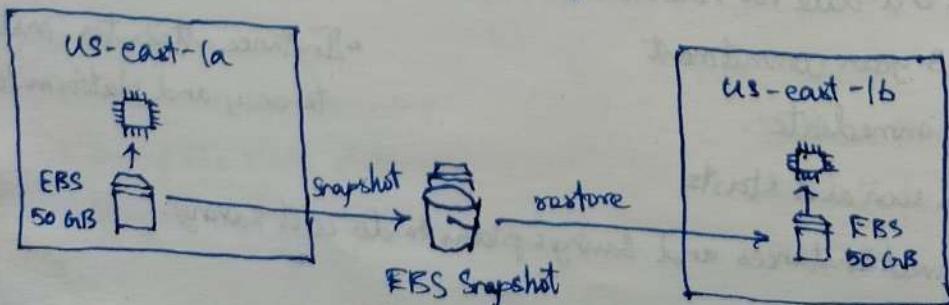
- An Elastic Block Store volume is a network drive you can attach to your instances.
- It allows your instances to persist data, even after their termination.
- They can only be mounted to one instance at a time. Multi-attach for some EBS
- They are bound to a specific availability zone.
- Analogy: Think of them as a 'Network USB stick'
- Free Tier: 30 GB of free EBS storage of General Purpose (SSD) or Magnetic per month
- It is a network drive (i.e. not a physical drive). It uses the network to communicate to the instance, which means there might be a bit of latency
- It can be detached from an EC2 instance and attached to another one quickly.
- To move a volume across, you first need to snapshot it.
- Have a provisioned capacity (size in GiBs and IOPS)



- Delete on termination attribute controls the EBS behaviour when an EC2 instance terminates.
- By default root volume is deleted. Other EBS volumes attribute is not marked.
- Use case: preserve root/EBS volume when instance is terminated.

## EBS Snapshots

- Make a backup (snapshot) of EBS volume at a point in time.
- Not necessary to detach volume to do snapshot but recommended.
- Can copy snapshots across AZ or region

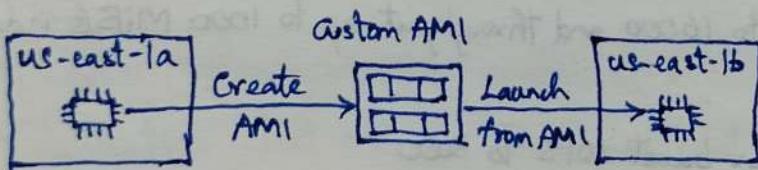


## AMI Overview

- Amazon Machine Image is customization of EC2 instance
- You add your own software, configuration, operating system, monitoring...
- Faster boot/configuration time because all your software is pre-packed
- AMI are built for a specific region (and can be copied across regions)
- You can launch EC2 instances from:
  - A public AMI : AWS provided
  - Your own AMI : You make and maintain them yourself
  - AWS Marketplace AMI : an AMI someone else made (and potentially sells)

## AMI Process [from EC2 instance]

- Start an EC2 instance and customize it
- Stop the instance (for data integrity)
- Build an AMI - this will also create EBS snapshots
- Launch instances from other AMIs



## EC2 Instance Store

- EBS volumes are network drives with good but "limited" performance.
- If you need a high-performance hardware disk, use EC2 Instance Store.
- Better I/O performance
- EC2 instance store lose their storage if they're stopped (ephemeral)
- Good for buffer/cache/scratch data/temporary content
- Risk of data loss if hardware fails
- Backups and Replication are your responsibility

## EBS Volume Types

- i) gp2/gp3 (SSD) → General purpose SSD - balances price & performance for wide variety.
  - ii) io1/io2 (SSD) → Highest performance SSD, mission critical low-latency or high throughput
  - iii) st1 (HDD) → Low cost HDD volume designed for frequently accessed intensive workload
  - iv) sc1 (HDD) → Lowest cost HDD volume designed for less frequently accessed workloads
- EBS volumes are characterized in Size | Throughput | IOPS (I/o per sec)
  - Only gp2/gp3 and io1/io2 can be used as boot volumes.

### i) General Purpose SSD

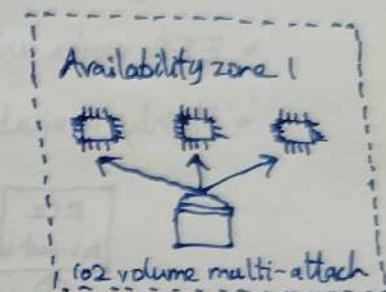
- Cost effective storage, low latency.
- System boot volumes, Virtual desktops, Development and test environments
- 1 GiB - 16 TiB
- gp3
  - Baseline of 3000 IOPS and throughput of 125 MiB/s
  - Can increase IOPS up to 16000 and throughput up to 1000 MiB/s independently
- gp2
  - Small gp2 volumes can burst IOPS to 3000
  - Size of the volume and IOPS are linked, max IOPS is 16,000
  - 3 IOPS per GiB, means at 53.34 GiB we are at max IOPS

### ii) Partitioned IOPS SSD

- Critical business applications with sustained IOPS performance or more than 16000 IOPS
- Great for databases, workloads (sensitive to storage perf and consistency)
- io1/io2 (4 GiB - 16 TiB)
  - Max PIOPS : 64000 for Nitro EC2 instances & 32000 for other
  - Can increase PIOPS independently from storage size
  - io2 have more durability and more IOPS per GiB (at the same price as io1)
- io2 Block Express (4 GiB - 64 TiB)
  - Sub-millisecond latency
  - Max PIOPS : 256000 with an IOPS:GiB ratio of 1000:1
- Supports EBS Multi-attach

## Hard Disk Drives (HDD)

- Cannot be a boot volume.
- 125 MiB to 16TiB
- Throughput Optimized HDD (st)
  - Big data, Data warehouses, Log processing
  - Max throughput 500 MiB/s - max IOPS 500
- Cold HDD (sc1)
  - For data that is infrequently accessed
  - Scenarios where lowest cost is important
  - Max throughput 250 MiB/s - max IOPS 250



## EBS Multi-Attach

- Exclusively for io1 / io2 family
- Attach the same EBS volume to multiple EC2 instances in the same AZ
- Each instance has full read & write permission to the volume
- Use case:
  - Achieve higher application availability in clustered Linux application (Eg: Teradata)
  - Application must manage concurrent write operations
- Must use a file system that's cluster aware (not XFS, EX4, etc...)

## EBS Encryption

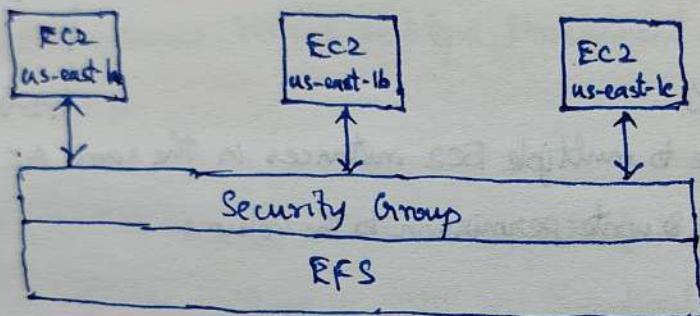
- Data at rest is encrypted inside the volume.
- All the data in flight moving between the instance and the volume is encrypted
- All snapshots are encrypted. All volumes created from the snapshot
- Encryption and decryption are handled transparently (you have nothing to do)
- Encryption has a minimal impact on latency.
- EBS Encryption leverages keys from KMS (AES-256)
- Copying an unencrypted snapshot allows encryption
- Snapshots of encrypted volumes are encrypted

## Encrypt an unencrypted EBS volume

- Create an EBS snapshot of the volume.
- Encrypt the EBS snapshot (using copy)
- Create new EBS volume from the snapshot (the volume will also be encrypted)
- Now you can attach the encrypted volume to the original instance

## EFS - Elastic File System

- Managed NFS (network file system) that can be mounted on many EC2 instances
- EFS works with EC2 instances in multi-AZ
- Highly available, scalable, expensive ( $3 \times gp2$ ) • Pay per use



- Use cases: content management, web serving, data sharing, WordPress
- Uses NFSv4.1 protocol
- Uses security group to control access to EFS
- Compatible with Linux based AMI (not Windows)
- Encryption at rest using KMS
- POSIX file system (~Linux) that has a standard file API
- File system scales automatically • pay per use, no capacity planning

## EFS - Performance & Storage Classes

### a) EFS Scale

- 1000s of concurrent NFS clients, 10 GiB/s throughput
- Grow to Petabyte-scale network file system, automatically.

## b) Performance mode (set at EFS creation time)

- General purpose (default): latency-sensitive use cases (web servers, CMS, etc...)
- Max I/O - higher latency, throughput, highly parallel (big data, media processing)

## c) Throughput mode

- Bursting (1TB = 50 MiB/s + burst of up to 100 MiB/s)
- Provisioned: set your throughput regardless of storage size, ex: 1 GiB/s for 1TB storage

## d) Storage Tiers (Lifecycle management feature - move file after N days)

- Standard: for frequently accessed files
- Infrequent access (EFS-IA): cost to retrieve files, lower price to store

### EBS

vs

### EFS

- Can be attached to one instance at a time
- Looked at the Availability zone (AZ) level
- gp2: IO increases if the disk size increases
- io1: can increase IO independently

To migrate EBS volume across AZ

- Take a snapshot
- Restore the snapshot to another AZ
- EBS backups use IO and you shouldn't run them which application is handling lot of traffic
- Root EBS volumes of instances get terminated by default if the EC2 instance gets terminated (you can disable that)

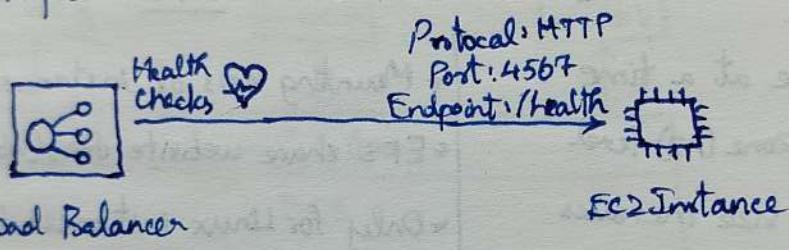
- Mounting 100s of instances across AZ
- EFS share website files (Wordpress)
- Only for Linux instances (POSIX)
- EFS has a higher price point than EBS
- Can leverage EFS-IA for cost savings

## ELB - Elastic Load Balancing

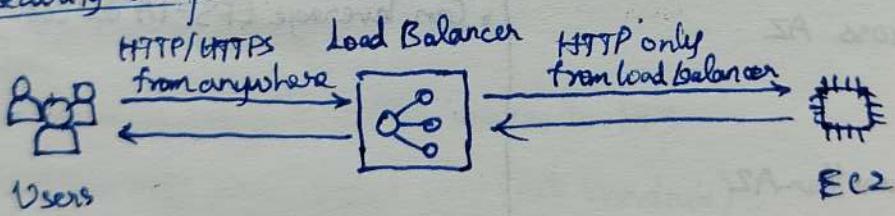
- An Elastic Load Balancer is a managed load balancer
- AWS takes care of upgrades, maintenance, high availability
- It costs less to setup your own load balancer but it will be a lot more effort on your end
- It is integrated with many AWS offerings/services

## Health Checks

- Health checks are crucial for Load Balancers
- They enable the load balancer to know if instances it forwards traffic to are available to reply to requests.
- The health check is done on a port and a route (/health is common)
- If the response is not 200 (OK), then the instance is unhealthy

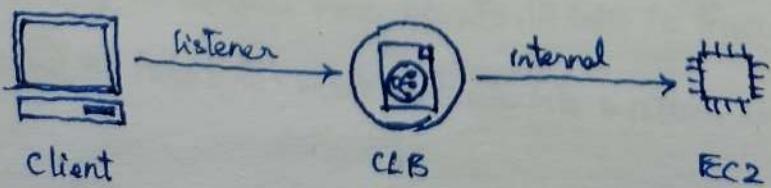


## Load Balancer Security Groups



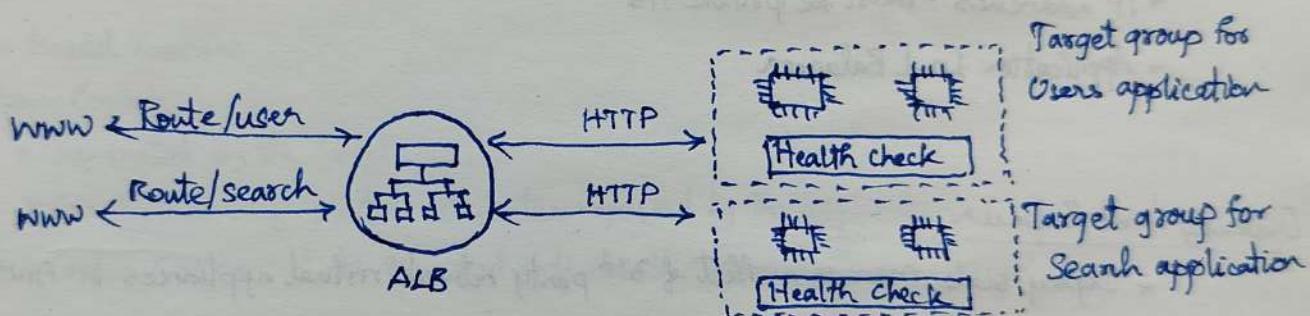
## Classic Load Balancers

- Supports TCP (Layer 4) - HTTP & HTTPS (Layer 7)
- Health checks are TCP or HTTP based
- Fixed hostname `XX.XX.region.elb.amazonaws.com`



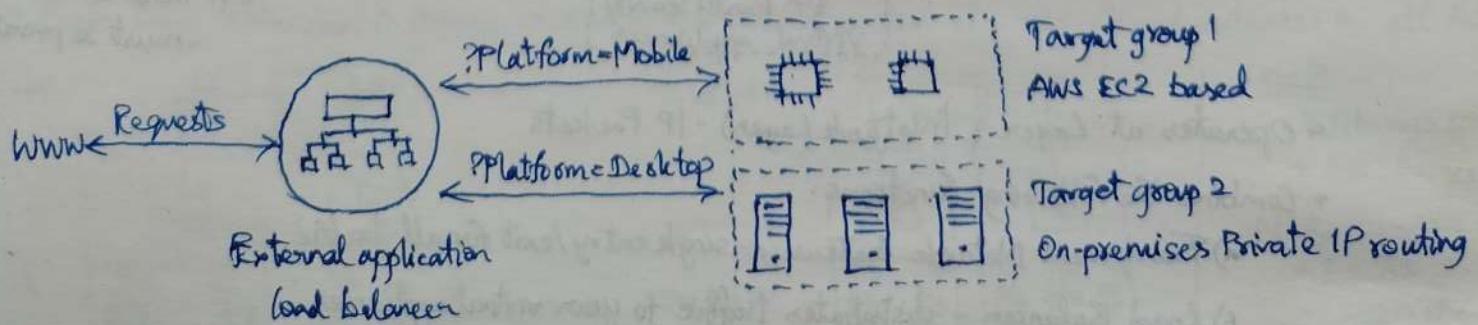
## Application Load Balancer

- Application Load Balancer is Layer 7 (HTTP)
- Load balancing to multiple HTTP applications across machines (target groups)
- Load balancing to multiple applications on the same machine (Eg. containers)
- Support for HTTP/2 and WebSocket
- Support redirects (from HTTP to HTTPS for example)
- Routing tables to different target groups:
  - Routing based on path in URL (example.com/users & example.com/posts)
  - Routing based on hostnames in URL (one.example.com & other.example.com)
  - Routing based on Query String / Headers (example.com/users?id=123&order=false)
- ALB are a great fit for micro services & container-based application (Eg. Docker, ECS)
- Has a port mapping feature to redirect to a dynamic port in ECS



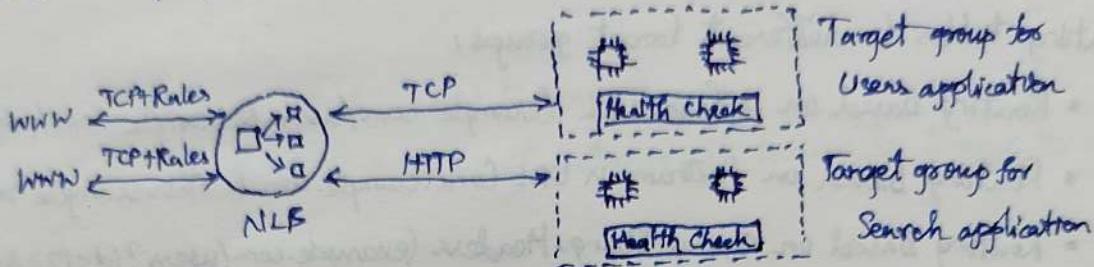
## Target Groups - ALB

- EC2 instances (can be managed by an Auto Scaling Group) - HTTP
- EBS tasks (managed by ECS itself) - HTTP
- Lambda functions - HTTP request is translated into a JSON event
- IP addresses - must be private IPs
- ALB can route to multiple target groups.
- Health checks are at the target group level



## Network Load Balancer

- It is Layer 4 that handle millions of requests per second
- Forward TCP & UDP traffic to your instances.
- Less latency ~100 ms (vs 400 ms for ALB)
- It has one static IP per AZ and supports assigning Elastic IP (helpful for whitelisting specific IP)
- Used for extreme performance TCP or UDP traffic

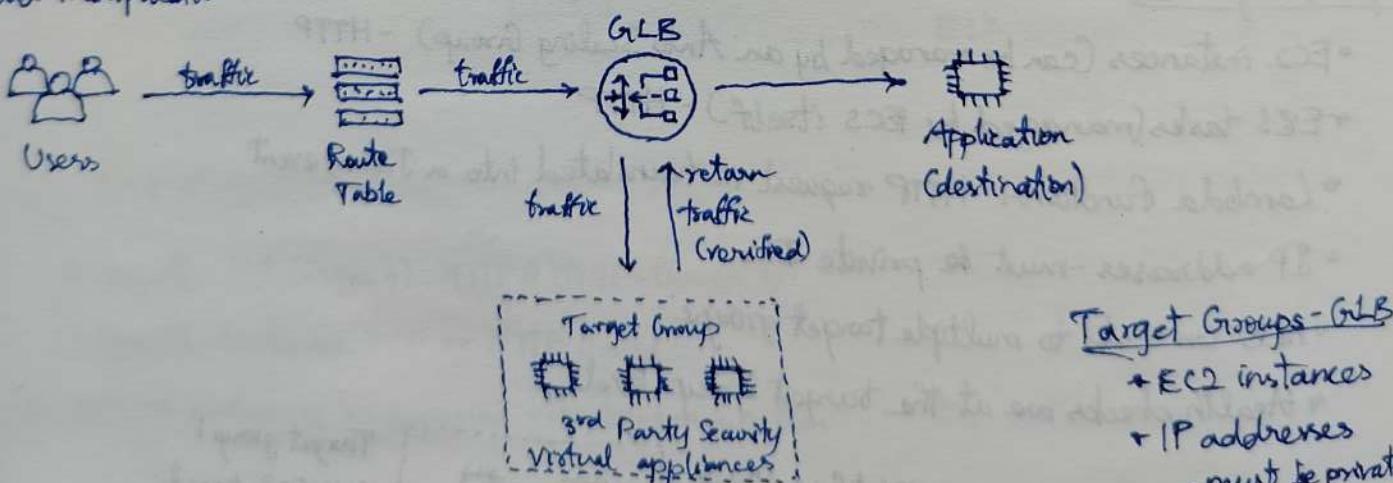


## Target Groups - NLB

- EC2 instances
- IP addresses - must be private IPs
- Application Load Balancer

## Gateway Load Balancer

- Deploy, scale, manage a fleet of 3<sup>rd</sup> party network virtual appliances in AWS
- Eg: Firewalls, Intrusion Detection and Prevention Systems, Deep Packet Inspection systems, payload manipulation.



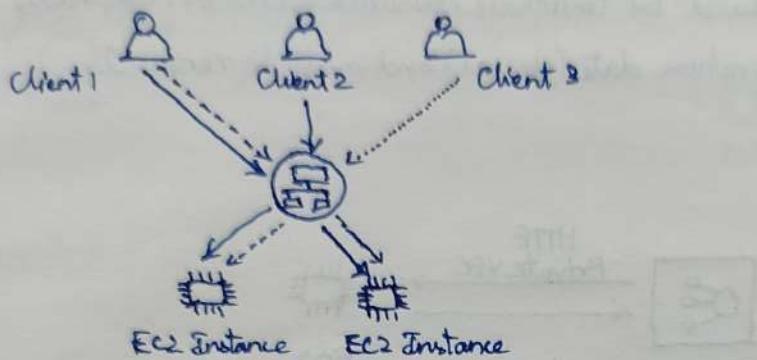
## Target Groups - GLB

- EC2 instances
- IP addresses
- must be private IPs

- Operates at Layer 3 (Network Layer) - IP Packets
- Combines the following functions:
  - a) Transparent Network Gateway - single entry/exit for all traffic
  - b) Load Balancer - distributes traffic to your virtual appliances
- Uses the GENEVE protocol on port 6081

## Sticky Sessions (Session Affinity)

- Implement stickiness so that the same client is always redirected to the same instance behind a load balancer. It works for Classic Load Balancer & Application Load Balancer
- The "cookie" used for stickiness has an expiration date you control.
- Use case: make sure the user doesn't lose his session data
- Enabling stickiness may bring imbalance to the load over the backend EC2 instances



## Cookie Names

### i) Application based Cookies

#### a) Custom Cookie

- Generated by the target
- Can include any custom attributes required by the application
- Cookie name must be specified individually for each target group

#### b) Application cookie

- Generated by the load balancer
- Cookie name is AWSALBAPP

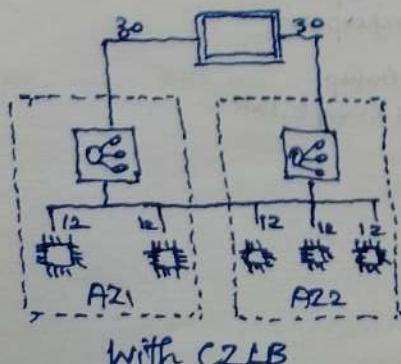
### ii) Duration based Cookies

- Cookie generated by the load balancer
- Cookie name is AWSALB for ALB, AWSELB for CLB

## Cross-Zone Load Balancing

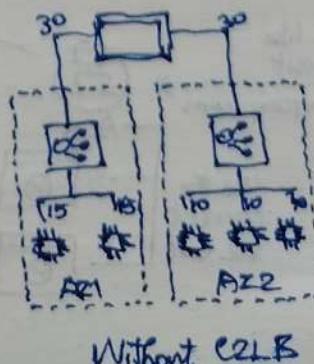
- Each load balancer instance distributes evenly across all registered instances in all AZ

Total - 60



Total - 60

AZ1 - 15  
AZ2 - 10



ALB

- Always ON
- No charges

NLB

- Disabled
- Charges if enable

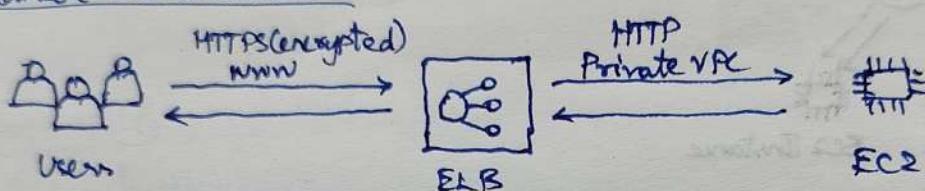
CLB

- Disabled
- Non-charger

## SSL/TLS

- SSL certificate allows traffic between your clients and your load balancer to be encrypted in transit (in-flight encryption)
- SSL → Secure Sockets Layer used to encrypt connections
- TLS → Transport Layer Security newer version
- Nowadays TLS certificates are mainly used but people still refer as SSL
- Public SSL certificates are issued by Certificate Authorities (CA) Eg: GoDaddy
- SSL certificates have an expiration date (you set) and must be renewed.

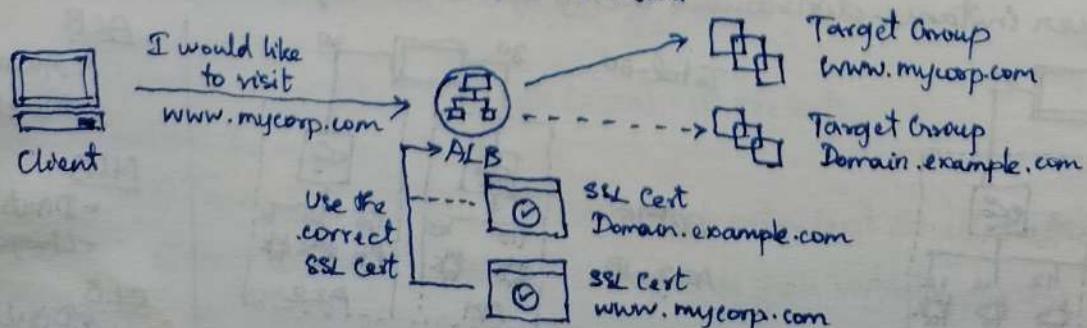
## Load Balancer - SSL Certificates



- Load Balancer uses an X.509 certificate (SSL/TLS server certificate)
- You can manage certificates using ACM (AWS Certificate Manager)
- You can create/upload your own certificates alternatively
- HTTPS Listener
  - You must specify a default certificate
  - You can add an optional list of certs to support multiple domains
  - Clients can use SNI (Server Name Indication) to specify the hostname they reach
  - Ability to specify a security policy to support older versions of SSL/TLS (legacy clients)

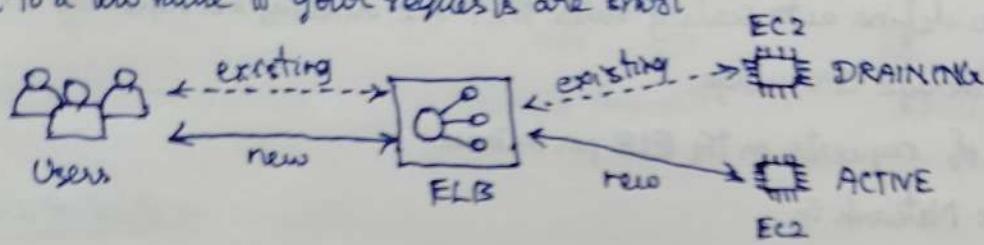
## Server Name Indication

- Solves the problem of loading multiple SSL certificates onto one web server (to serve multiple websites)
  - It's a newer protocol and requires the client to indicate the hostname of the target server in the initial SSL handshake. Server will find correct certificate or return the default one.
- Only works for ALB & NLB, CloudFront



## Connection Draining

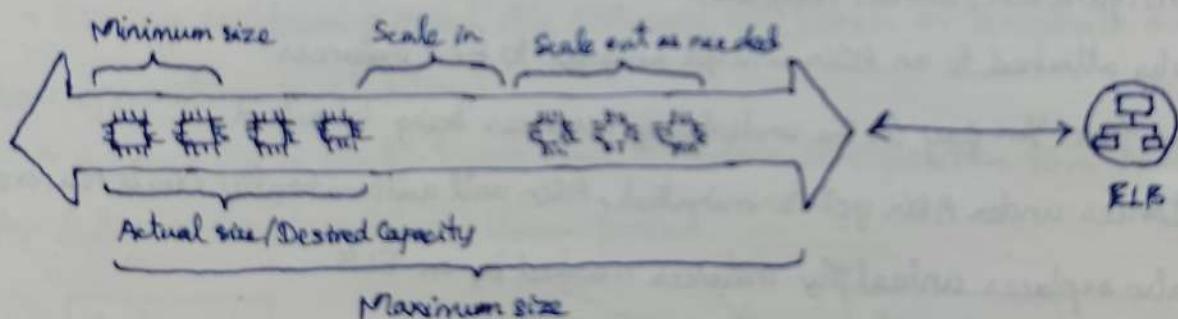
- Also known as 'Deregistration Delay'
- Time to complete 'in-flight' requests while the instance is de-registering or unhealthy
- Stops sending new requests to the EC2 instance which is de-registering.
- Between 1 to 3600 seconds (default: 300 seconds). It can be disabled (set value to 0)
- Set to a low value if your requests are short



## Auto-Scaling Group

Why?

- Scale out (add EC2 instances) to match an increased load
- Scale in (remove EC2 instances) to match a decreased load
- Ensure we have a minimum and maximum number of machines running.
- Automatically register new instances to a load balancer.



## ASG Attributes

- Launch configuration
  - AMI + Instance Type
  - EC2 User Data
  - EBS Volumes
  - Security Groups
  - SSTH Key Pair
- Min Size / Max Size / Initial Capacity
- Network + Subnets information
- Load Balancer Information
- Scaling policies

## Auto Scaling Alarms

- It is possible to scale an ASG based on CloudWatch alarms.
- An alarm monitors a metric (such as Average CPU)
- Metrics are computed for the overall ASG instances
- Based on the alarm we can create scale-out or scale-in policies
- It is possible to define auto scaling rules that are directly managed by EC2
  - Target Average CPU usage
  - Number of requests on the ELB per instance
  - Average Network in
  - Average Network out
- We can auto scale based on a custom metric (ex: no. of connected users)
  - Send custom metric from application on EC2 to CloudWatch (PutMetric API)
  - Create CloudWatch alarm to react to low/high values.
  - Use the CloudWatch alarm as the scaling policy for ASG
- ASGs use Launch Configuration or Launch Templates. To update an ASG, you must provide a new launch configuration/launch template.
  - IAM roles attached to an ASG will get assigned to EC2 instances.
  - ASGs are free. You pay for the underlying resources being launched
  - If instances under ASG get terminated, ASG will automatically create new ones as replacement. It also replaces unhealthy instances marked by an ELB

## Dynamic Scaling Policies

### a) Target Tracking Scaling

- Most simple and easy to set-up
- Eg: Average ASG CPU to stay at around 40%

### b) Simple/Step Scaling

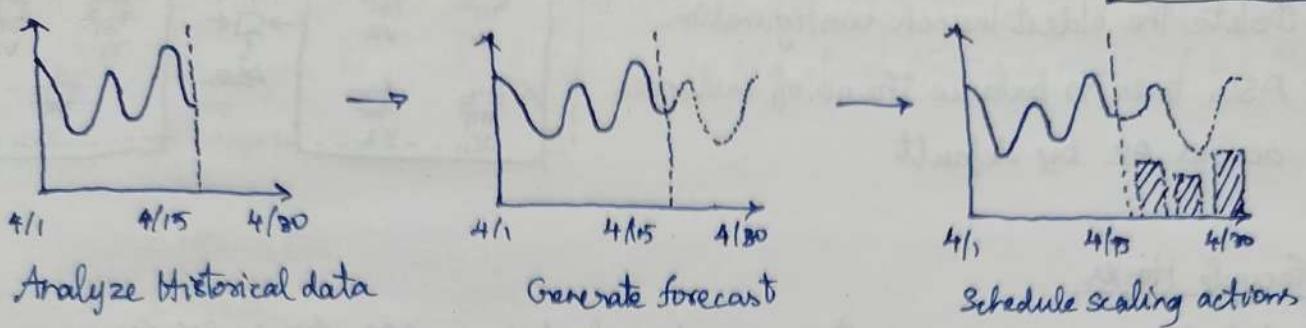
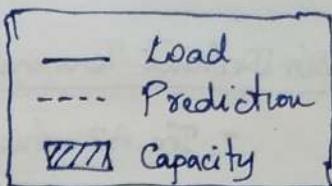
- CloudWatch alarm ( $CPU > 70\%$ ) then add 2 units
- CloudWatch alarm ( $CPU < 30\%$ ) then remove 1 unit

### c) Scheduled Actions

- Anticipate a scaling based on known usage patterns
- Eg: Increase the min capacity to 10 at 5 pm on Fridays

## Predictive Scaling

- Continuously forecast load and schedule scaling ahead.

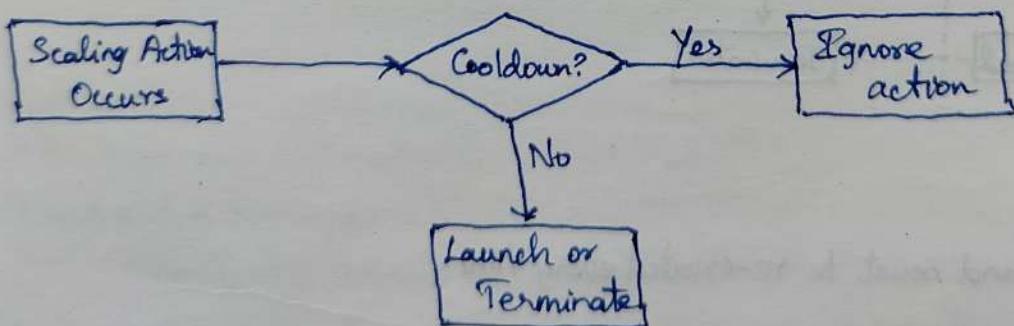


## Good Metrics to Scale on

- CPU Utilization : Average CPU utilization across your instances
- RequestCountPerTarget : Make sure no. of requests per EC2 instance is stable
- Average Network In/Out : If your application is network bound
- Any custom metric : Push using CloudWatch

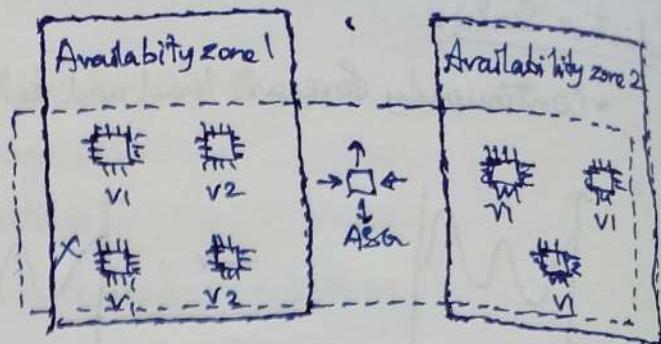
## Scaling Cooldowns

- After a scaling activity happens, it will be in cooldown period (default 300 seconds)
- During cooldown period, the ASG will not launch or terminate additional instances (to allow for metrics to stabilize)
- Advice: Use a ready-to-use AMI to reduce configuration time in order to serve requests faster and reduce the cooldown period.



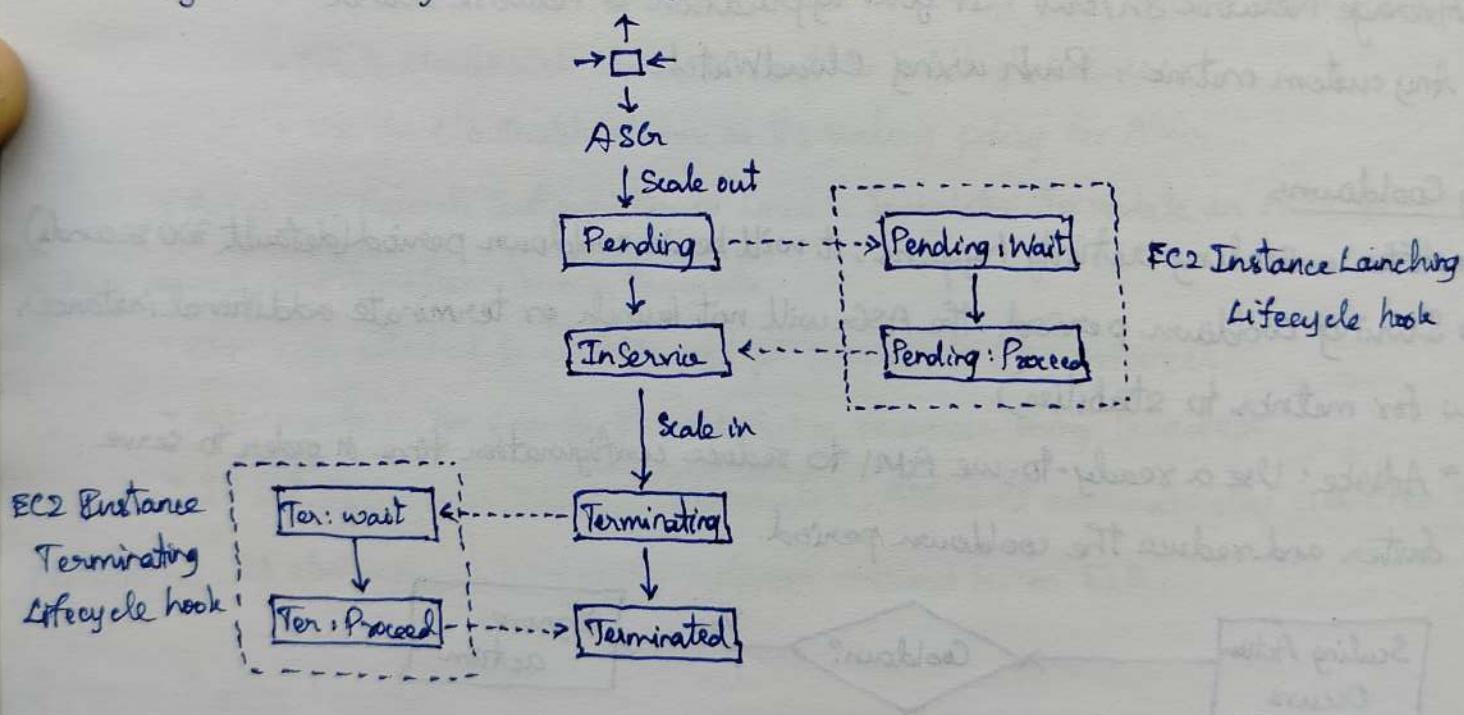
## ASG Default Termination Policy

- The AZ which has most no. of instances
- Delete the oldest launch configuration
- ASG tries to balance the no. of instances across AZ by default



## ASG Lifecycle Hooks

- By default as soon as an instance is launched in an ASG it's in service.
- You have the ability to perform extra steps before the instance goes in service (Pending state before the instance is terminated (Terminating state))
- Eg: Extract logs, data or any other info before the instance is terminated.



## Launch Configuration

- It is legacy and must be re-created every time.

## Launch Template

- Can have multiple versions
- Create parameters subsets (partial configuration for re-use and inheritance)
- Provision using both On-demand and Spot instances (or a mix)
- Can use T2 unlimited burst feature

## AWS RDS Overview

- RDS stands for Relational Database Service
- It's a managed DB service for DB use SQL as a query language
- It allows you to create databases in the cloud that are managed by AWS
  - Postgres
  - MySQL
  - MariaDB
  - Oracle
  - Microsoft SQL Server
  - Aurora

## Advantages over DB on EC2

- It is a managed service.
- Automated provisioning, OS patching, Monitoring dashboards
- Continuous backups and restore to specific timestamp (Point in Time Restore)
- Read replicas for improved read performance
- Multi AZ setup for DR (Disaster Recovery)
- Maintenance windows for upgrades
- Scaling capability (vertical and horizontal)
- Storage backed by EBS
- But you can't SSH into the RDS instances.

## RDS Backups

- Backups are automatically enabled in RDS
- a) Automated backups:
- Daily full backup of the database (during the maintenance window)
  - Transaction logs are backed by RDS every 5 minutes
  - Ability to restore to any point in time (from oldest backup to 5 minutes ago)
  - 7 days retention (can be increased to 35 days)

b) DB Snapshots:

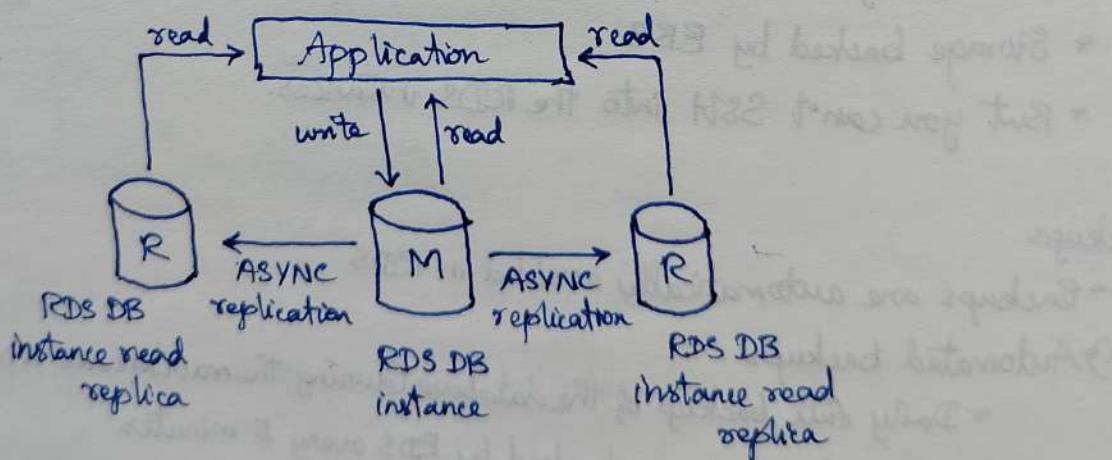
- Manually triggered by the user.
- Retention of backup for as long as you want.

## RDS - Storage Auto Scaling

- Helps you increase storage on your RDS DB instance dynamically.
- When RDS detects you are running out of free database storage, it scales automatically.
- Avoid manually scaling your database storage.
- You have to set Maximum Storage Threshold (maximum limit for DB storage)
- Automatically modify storage if:
  - Free storage is less than 10% of allocated storage
  - Low storage lasts at least 5 minutes
  - 6 hours have passed since last modification.
- Useful for applications with unpredictable workloads

## RDS Read Replicas

- Up to 5 Read Replica
- Within AZ, Cross AZ or Cross Region
- Replication is ASYNC, so reads are eventually consistent.
- Replicas can be promoted to their own DB
- Applications must update the connection string to leverage read replicas.



## Use Case:

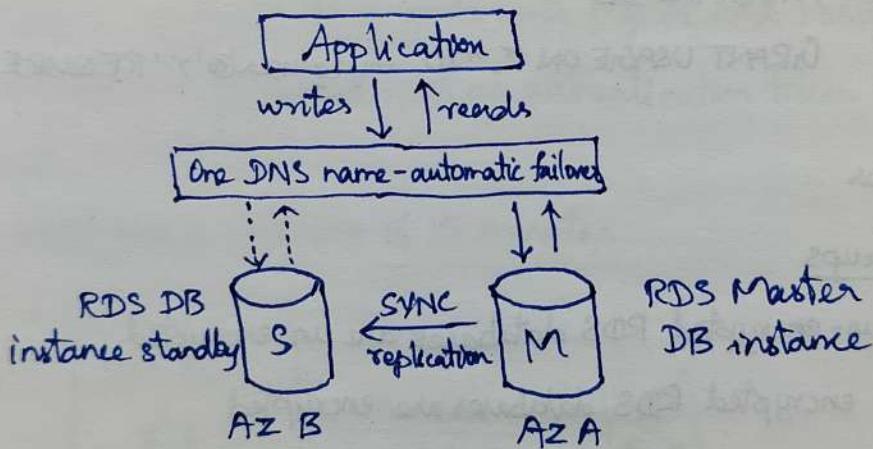
- You have a production database that is taking on normal load.
- You want to run a reporting application to run some analytics.
- You create a Read Replica to run the new workload so production application is unaffected.
- Read replicas are used for SELECT statement only
- Do not use INSERT, UPDATE, DELETE as it will modify the DB

## RDS - Network Cost

- In AWS, there is a network cost when data goes from one AZ to another
- For RDS Read replicas within the same region, you don't pay that fee.

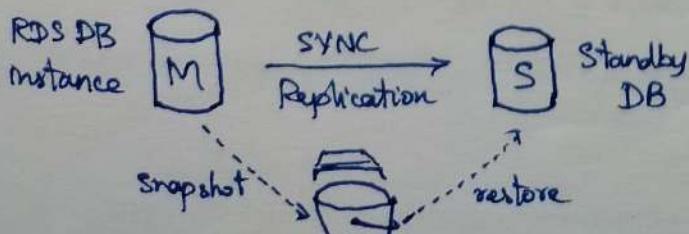
## RDS Multi AZ

- SYNC replication
- One DNS name - automatic app failover to standby.
- Increase availability.
- Failover in case of loss of AZ, loss of network, instance or storage failure.
- No manual intervention in apps.
- Not used for scaling.
- Note: Read Replicas can be setup as Multi AZ for Disaster Recovery.



## RDS Single AZ to Multi AZ

- Zero downtime operation (no need to stop the DB)
- Just click on "modify" for the database and enable multi AZ
- The following happens behind the scenes:
  - A snapshot is taken
  - New DB is restored from the snapshot in a new AZ
  - Synchronization is established between the two databases



## RDS Security - Encryption

### a) At rest encryption

- Possibility to encrypt the master and read replicas with AWS KMS - AES 256
- Encryption has to be defined at launch time
- If the master is not encrypted, the read replicas cannot be encrypted
- Transparent Data Encryption (TDE) available for Oracle and SQL Server

### b) In-flight encryption

- SSL certificate to encrypt data to RDS in flight
- Provide SSL options with trust certificate when connecting to database
- To enforce SSL:
  - PostgreSQL : rds-force-ssl=1 in the AWS RDS Console (Parameter Groups)
  - MySQL : Within the DB  
`GRANT USAGE ON *.* TO 'mysqluser@%'; REQUIRE SSL;`

## RDS Encryption Operations

### a) Encrypting RDS backups

- Snapshots of un-encrypted RDS databases are un-encrypted
- Snapshots of encrypted RDS databases are encrypted
- Can copy an unencrypted snapshot into an encrypted one

### b) Encrypting RDS database:

- Create a snapshot of the unencrypted database
- Copy the snapshot and enable encryption for the snapshot
- Restore the database from the encrypted snapshot
- Migrate applications to the new database and delete the old database

## RDS Security

### a) Network security

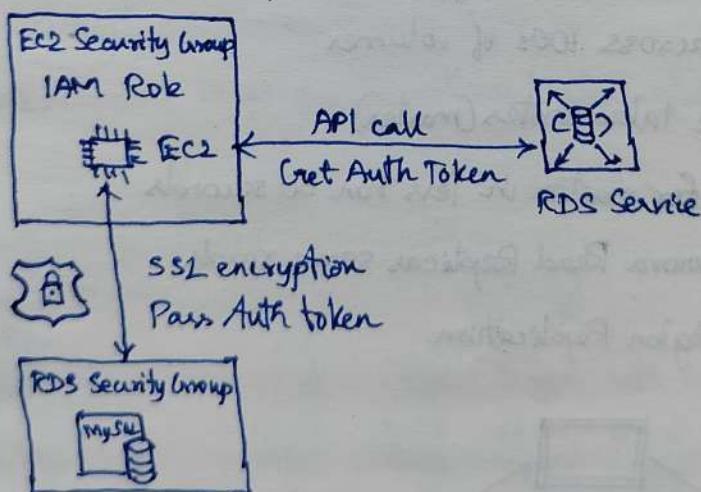
- RDS databases are usually deployed within a private subnet, not in a public one
- RDS security works by leveraging security groups - it controls which IP or security group can communicate with RDS

### b) Access Management

- IAM policies help control who can manage AWS RDS (through the RDS API)
- Traditional username and password can be used to login into the database.
- IAM based authentication can be used to login into RDS MySQL & PostgreSQL

## RDS IAM Authentication

- IAM database authentication works with MySQL and PostgreSQL
- You don't need a password, just an authentication token obtained through IAM & RDS API calls
  - Auth token has a lifetime of 15 minutes



## Benefits

- Network in/out must be encrypted using SSL
- IAM to centrally manage users instead of DB
- Can leverage IAM roles and EC2 instance profiles for easy integration.

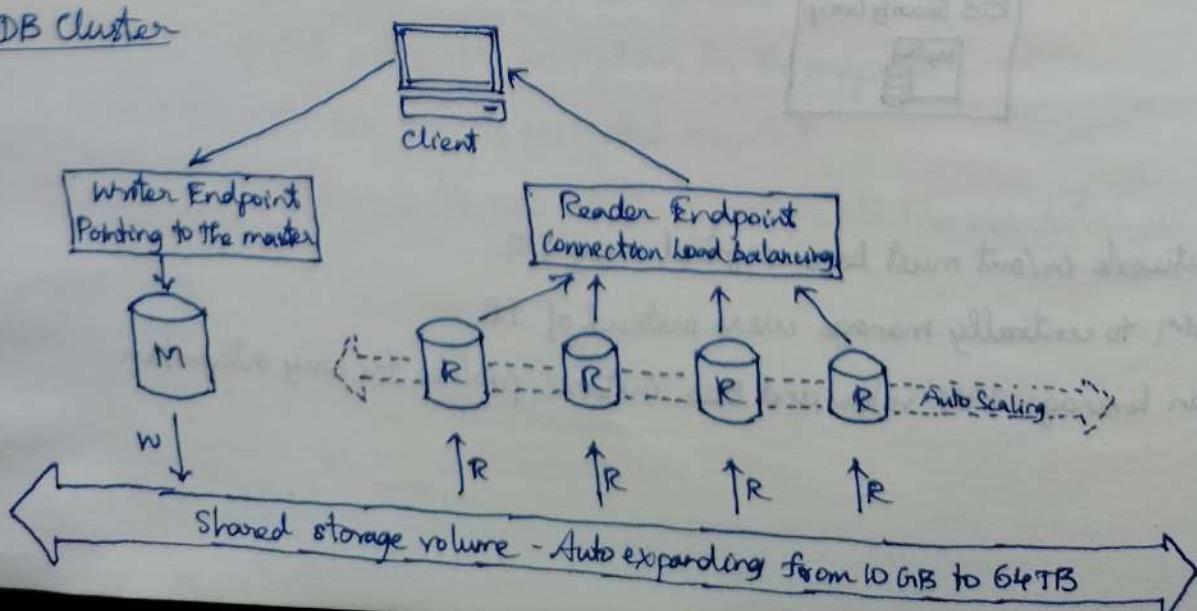
## Amazon Aurora

- Proprietary technology from AWS (not open sourced)
- PostgreSQL and MySQL are both supported as AuroraDB
- It is cloud-optimized. 5x performance improvement over MySQL on RDS, over 3x the performance of PostgreSQL on RDS
- Aurora storage automatically grows in increments of 10 GB up to 128 TB
- Aurora can have 15 replicas (MySQL has 5) and the replication process is faster (sub 10 ms replica lag)
- Failover in Aurora is instantaneous. It's High availability native
- Aurora costs more than RDS (20% more) - but it is more efficient.

## Aurora - High Availability

- 6 copies of your data across 3 AZ
  - 4 copies out of 6 needed for writes
  - 3 copies out of 6 needed for reads
- Self-healing with peer-to-peer replication
- Storage is striped across 100s of volumes
- One Aurora Instance takes writes (master)
- Automated failover for master in less than 30 seconds
- Master + up to 15 Aurora Read Replicas serve reads
- Support for Cross Region Replication

## Aurora DB Cluster



## Aurora - Features

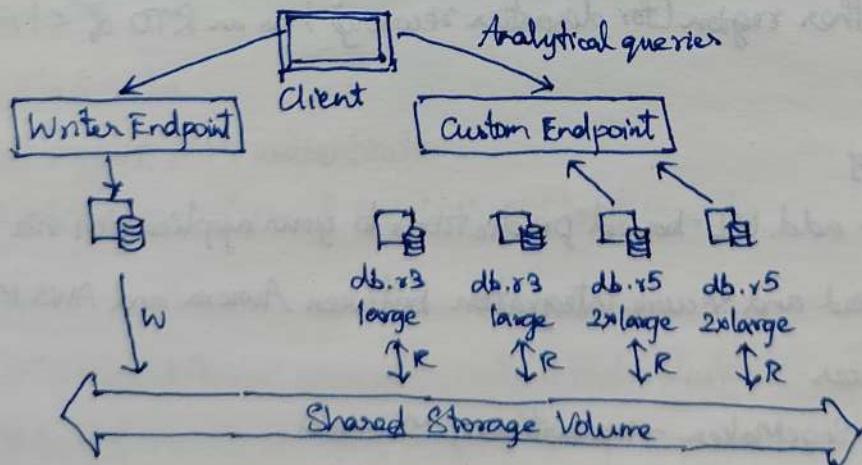
- Automatic fail-over
- Backup and Recovery
- Isolation and security
- Industry compliance
- Backtrack : restore data at any point of time without using backups
- Push-button scaling
- Automated patching with zero downtime
- Advanced Monitoring
- Routine Maintenance

## Aurora - Security

- Similar to RDS because uses the same engines.
- Encryption at-rest using KMS
- Automated backups, snapshots and replicas are also encrypted
- Encryption in flight using SSL
- Possibility to authenticate using IAM token
- You are responsible for protecting the instance with security groups. You can't SSH

## Aurora - Custom Endpoints

- Define a subset of Aurora instances as a Custom Endpoint
- Eg : Run analytical queries on specific replicas
- The Reader Endpoint is generally not used after defining Custom Endpoints

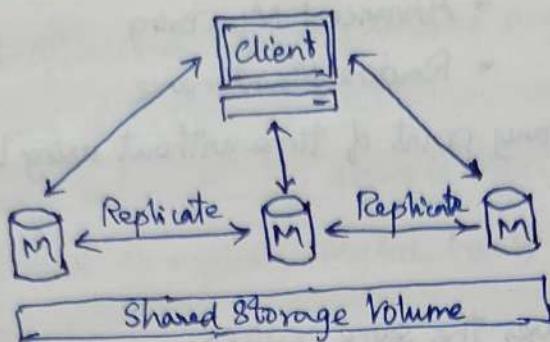


## Aurora - Serverless

- Automated database instantiation and auto scaling based on actual usage.
- Good for infrequent, intermittent or unpredictable workloads
- No capacity planning needed. Pay per second can be more cost-effective
- Client is going to talk to 'Proxy fleet' many Aurora instances will be created

## Aurora - Multi Master

- For immediate failover of write node to achieve high availability
- Every node does R/W - vs promoting a RR as the new master



## Global Aurora

### a) Aurora Cross Region Read Replicas

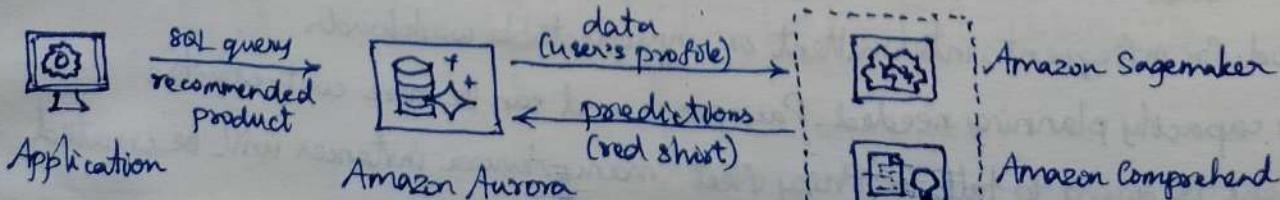
- Useful for disaster recovery
- Simple to put in place

### b) Aurora Global Database

- One primary region (read/write)
- Up to 5 secondary (read-only) regions, replication lag is less than 1 second
- Up to 16 Read Replicas per secondary region
- Helps for decreasing latency
- Promoting another region (for disaster recovery) has an RTO of <1 minute

## Aurora - Machine Learning

- Enables you to add ML-based predictions to your application via SQL
- Simple, optimized and secure integration between Aurora and AWS ML services
- Supported services
  - Amazon SageMaker - use with any ML model
  - Amazon Comprehend - for sentiment analysis
- Use cases: Fraud detection, ads targeting, product recommendations

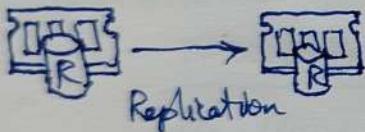


## Amazon ElastiCache

- To get managed Redis or Memcached
- Caches are in-memory databases with really high performance, low latency
- Helps reduce load off of databases for read intensive workloads
- Helps make your application stateless
- AWS takes care of OS maintenance / patching, optimization, setup, configuration monitoring, failure recovery and backups.
- Using ElastiCache involves heavy application code changes.

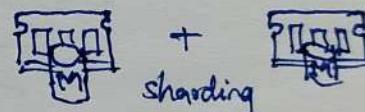
### REDIS

- Multi AZ with Auto-Failover
- Read Replicas to scale reads and have high availability
- Data durability using AOF persistence
- Backup and restore features



### MEMCACHED

- Multi-node for partitioning of data (sharding)
- No high availability (replication)
- Non persistent
- Multi-threaded architecture
- No backup and restore



## ElastiCache - Security

### a) All Cache in ElastiCache

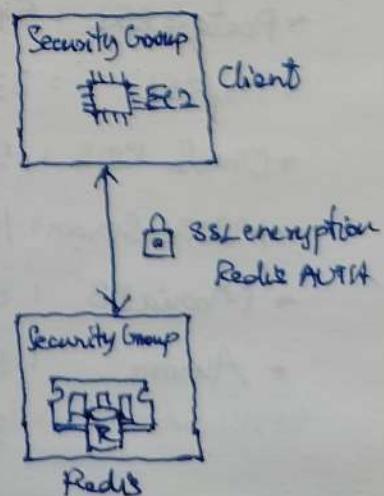
- Do not support IAM authentication
- IAM policies are only used for AWS-API level security

### b) Redis AUTH

- Set a 'password/token' when you create a Redis cluster
- Extra level of security on top of security groups
- Support SSL in flight encryption

### c) Memcached

- Supports SASL-based authentication (advanced)



## ElasticCache - Patterns

- Lazy Loading : All the read data is cached. Data can become stale in cache
- Write Through : Adds or update data in the cache when written to a DB (no stale data)
- Session Store : Store temporary session data in a cache (using TTL features)

## ElasticCache - Redis Use Case

- Gaming leaderboards are computationally complex
- Redis sorted sets guarantee both uniqueness and element ordering (1, 2, 3...)
- Each time a new element is added, it's ranked in real time, then added in correct

## Important ports

- FTP : 21
- SSH : 22
- SFTP : 22
- HTTP : 80
- HTTPS : 443

## RDS Database ports

- PostgreSQL : 5432
- MySQL : 3306
- Oracle RDS : 1521
- MSSQL Server : 1433
- MariaDB : 3306
- Aurora : 5432 (if PostgreSQL)  
3306 (if MySQL)

## DNS

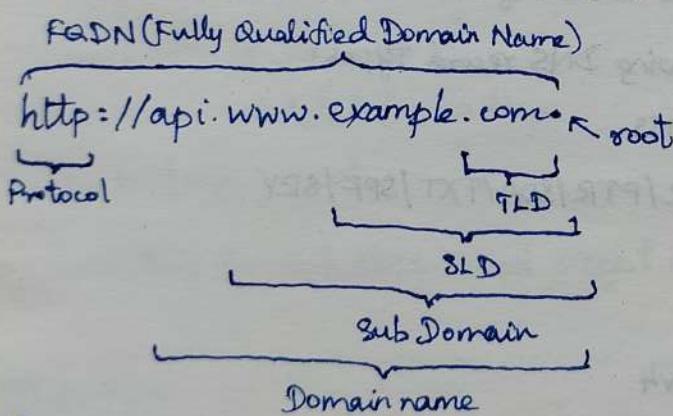
• Domain Name System which translates the human friendly hostnames into the machine IP addresses. Eg: www.google.com → 172.217.18.36

• DNS is the backbone of the Internet. It uses hierarchical naming structure

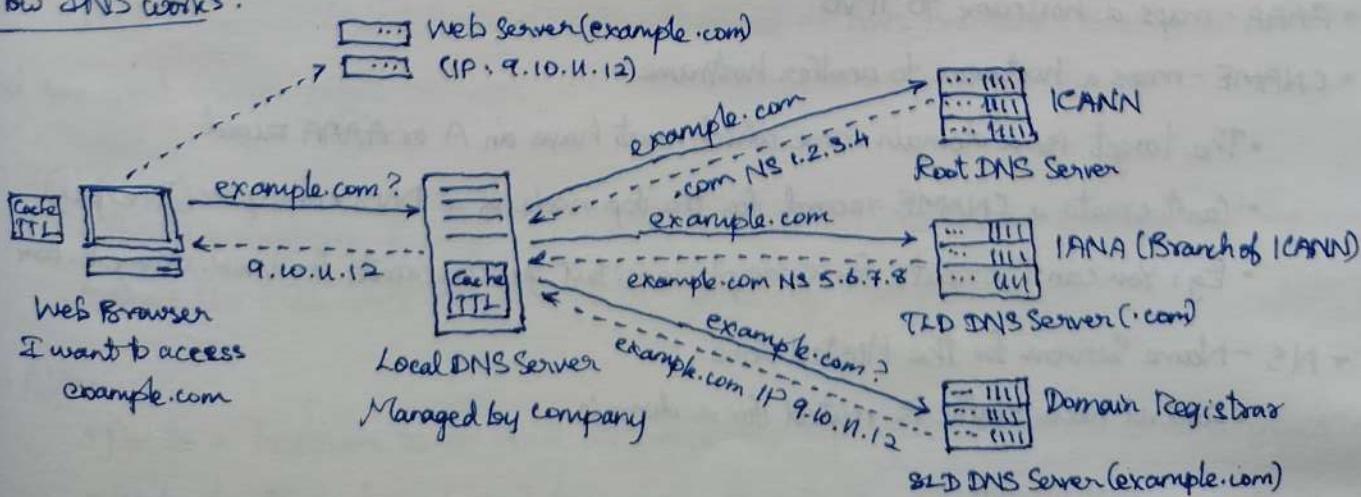
.com  
example.com  
www.example.com  
api.example.com

## DNS - Terminologies

- Domain Registrar : Amazon Route 53, GoDaddy
- DNS Records : A, AAAA, CNAME, NS
- Zone File : Contains DNS records
- Name Server : Resolves DNS queries (Authoritative or Non-Authoritative)
- Top Level Domain (TLD) : .com, .us, .in, .gov, .org
- Second Level Domain (SLD) : amazon.us.google.com



## How DNS works?



## Amazon Route 53

- Highly available, scalable, fully managed and Authoritative DNS
  - Authoritative → customer (you) can update the DNS records
- Route 53 is also a Domain Registrar
- Ability to check the health of your resources
- Only AWS service which provides 100% availability SLA
- Why number 53? 53 is a reference to the traditional DNS port

## Route 53 - Records

- Define how you want to route traffic for a domain
- Each record contains:
  - Domain/Subdomain name - Eg: example.com
  - Record Type. Eg: A or AAAA
  - Value. Eg: 12.34.56.78
  - Routing Policy - how Route 53 responds to queries
  - TTL - Time to live, the amount of time the record is cached at DNS Resolvers
- Route 53 supports the following DNS record types:
  - A/AAAA/CNAME/NS
  - CAA/DS/MX/NAPTR/PTR/SOA/TXT/SPF/SRV

## Route 53 - Record Types

- A - maps a hostname to IPv4
- AAAA - maps a hostname to IPv6
- CNAME - maps a hostname to another hostname
  - The target is a domain name which must have an A or AAAA record
  - Can't create a CNAME record for the top node of a DNS namespace (Zone apex)
  - Eg: You can't create for example.com but you can create for www.example.com
- NS - Name Servers for the Hosted Zone
  - Control how traffic is routed for a domain

## Route 53 - Hosted Zones

- A container for records that define how to route traffic to a domain and its subdomains

### a) Public Hosted Zones

- Contains records that specify how to route traffic on the internet (public domain names)  
application.mypublicdomain.com

### b) Private Hosted Zones

- Contains records that specify how you route traffic within one or more VPCs  
(private domain names) application.company.internal

- You pay \$0.50 per month per hosted zone, i.e. minimum of \$12 per year

## Route 53 - Records TTL (Time to Live)

### a) High TTL - eg: 24 hours

- Less traffic on Route 53
- Possibly outdated records

### b) Low TTL - eg: 60 seconds

- More traffic on Route 53 (\$)
- Records are outdated for less time
- Easy to change records

→ TTL is mandatory for each DNS record except for Alias records

## Route 53 - CNAME vs Alias

AWS Resources (Load Balancer, CloudFront) expose an AWS hostname: lb-1234.us-east2.elb.amazonaws.com

and you want myapp.mydomain.com

### a) CNAME

- Points a hostname to any other hostname (app.mydomain.com → blabla.anything.com)
- Only for Non-root domain (Eg: something.mydomain.com)

### b) Alias

- Points a hostname to an AWS Resource (app.mydomain.com → blabla.amazonaws.com)
- Works for Root domain and Non-root domain (Eg: mydomain.com)
- Free of charge and Native Health check capability

## Route 53 - Alias Records

- Maps a hostname to an AWS resource
- An extension to DNS functionality
- Automatically recognizes changes in the resource's IP addresses
- Unlike CNAME it can be used for the top node of a DNS namespace (Zone Apex)

Eg. example.com

- Alias Record is always of type A/AAAA for AWS resources (IPv4/IPv6)
- You can't set the TTL, it is set automatically by Route 53

## Route 53 - Alias Records Targets

- Elastic Load Balancers
- CloudFront Distributions
- API Gateway
- Elastic Beanstalk environments
- S3 Websites
- VPC Interface endpoints
- Global Accelerator
- Route 53 record in the same hosted zone

You cannot set an ALIAS record  
for an EC2 DNS name

## Route 53 - Routing Policy

- Define how Route 53 responds to DNS queries
- DNS does not route any traffic, it only responds to the DNS queries.
- Route 53 supports the following Routing Policies:

- a) Simple
- b) Weighted
- c) Failover
- d) Latency based
- e) Geolocation
- f) Multi-value answer
- g) Geoproximity (using Route 53 traffic flow feature)

## Simple Routing Policy

- Typically route traffic to a single resource.
- Can specify multiple values in the same record.
- If multiple values are returned, a random one is chosen by the client.
- When Alias is enabled, specifies only one AWS resource.
- Can't be associated with Health checks.

## Weighted Routing Policy

- Control the % of the requests that go to each specific resource.
- Assign each record a relative weight.

$$\text{Traffic \%} = \frac{\text{Weight for a specific record}}{\text{Sum of all weights for all records}}$$

Note: Weights don't need to sum up to 100.

- DNS records must have same name and type.
- Can be associated with Health checks.
- Use cases: Load balancing between regions, testing new application versions.
- Assign a weight of 0 to a record to stop sending traffic to a resource.

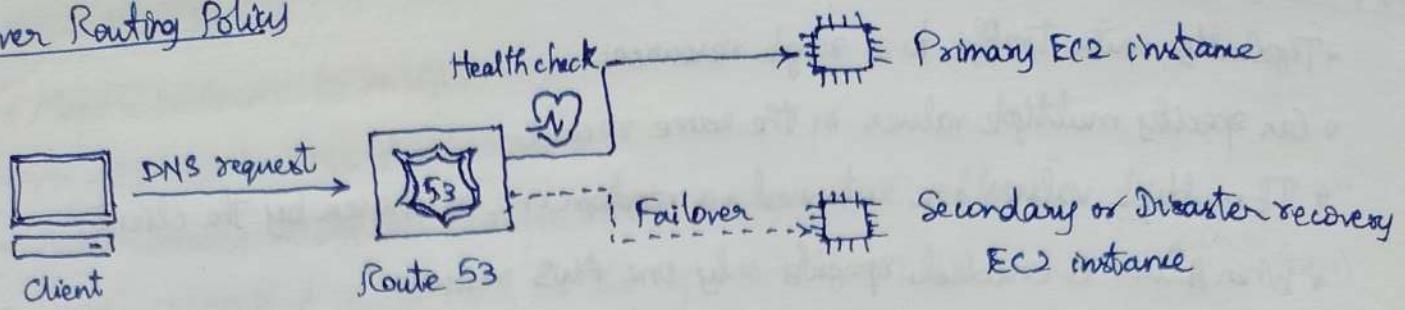
## Latency-based Routing Policy

- Redirect to the resource that has the least latency close to us.
- Super helpful when latency for users is a priority.
- Latency is based on traffic between users and AWS Regions.
- Can be associated with Health checks (has a failover capability).

## Geolocation Routing Policy

- Different from Latency-based. Routing is based on user location.
- Specify location by continent, country or by US state.
- Should create a "Default" record (in case there is no match on location).
- Can be associated with Health checks.
- Use cases: website localization, restrict content distribution, load balancing.

## Failover Routing Policy



## Multi-value Routing Policy

- Use when routing traffic to multiple resources.
- Route 53 returns multiple values/resources
- Can be associated with Health Checks (return only values for healthy resources)
- Up to 8 healthy records are returned for each Multi-value query
- Multi-value is not a substitute for having an ELB

## Geoproximity Routing Policy

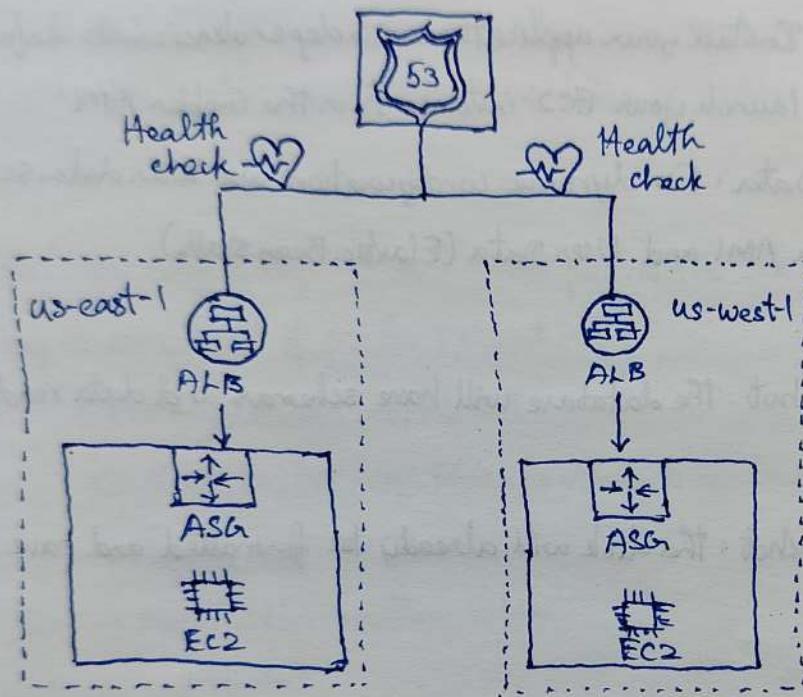
- Route traffic to your resources based on the geographic location of user and resource
- Ability to shift more traffic to resources based on the defined bias
- To change the size of the geographic region, specify bias values:
  - To expand (1 to 99) - more traffic to the resource
  - To shrink (-1 to -99) - less traffic to the resource
- Resources can be:
  - AWS resources (specify AWS region)
  - Non-AWS resources (specify Latitude and Longitude)
- You must use Route 53 Traffic flow (Advanced) to use this feature

## Route 53 - Traffic Flow

- Simplify the process of creating and maintaining records in large and complex configurations
- Visual editor to manage complex routing decision trees.
- Configurations can be saved as Traffic flow policy
  - Can be applied to different Route 53 Hosted Zones (different domain names)
  - Supports versioning

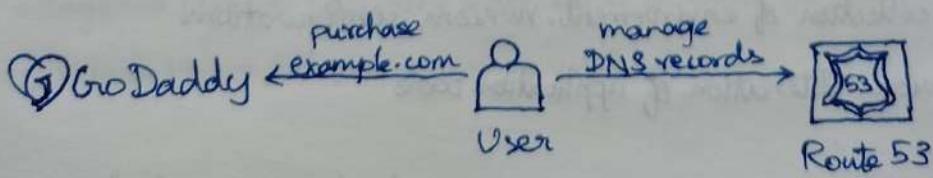
## Route 53 - Health Check

- \* HTTP Health check are only for public resources
- \* Automated DNS failover:
  - i) Health checks that monitor an endpoint - Application server, other AWS resources
  - ii) Health checks that monitor other health checks - Calculated Health Checks
  - iii) Health checks that monitor Cloudwatch Alarms - DynamoDB throttle alarms on RDS
- \* Health checks are integrated with CloudWatch metrics



## Domain Registrar vs DNS Service

- \* You buy or register your domain name with a Domain Registrar typically by paying annual charges. (Eg: GoDaddy, Amazon Registrar)
- \* The domain registrar usually provides you with a DNS Service to manage your DNS records
- \* But you can use another DNS service to manage your DNS records
- \* Eg: Purchase domain from GoDaddy and use Route 53 to manage your DNS records



Domain Registrar != DNS Service

## Instantiating Applications quickly

When launching a full stack (EC2, EBS, RDS) it can take time to:

- Install applications
- Insert initial (or recovery) data
- Configure everything
- Launch the application

### a) EC2 Instances

- Use a Golden AMI : Install your applications, OS dependencies etc beforehand and launch your EC2 instance from the Golden AMI
- Bootstrap using User Data : For dynamic configuration use User data scripts
- Hybrid : mix Golden AMI and User Data (Elastic Beanstalk)

### b) RDS Databases

- Restore from a snapshot : The database will have schemas and data ready

### c) EBS Volumes

- Restore from a snapshot : The disk will already be formatted and have data

## Elastic Beanstalk

- Developer centric view of deploying an application on AWS
- It is a managed service that handles EC2, ASG, ELB, RDS ...
- It automatically handles capacity provisioning, load balancing, scaling, monitoring, etc.
- Just the application code is responsibility of the developer
- We still have full control over the configuration
- Beanstalk is free but you pay for the underlying instances

## Components

- Application : collection of environment, versions, configurations
- Application Version : Iteration of application code
- Environment:
  - Collection of AWS resources running an application version (only one version at a time)
  - Tiers : Web Server Environment Tier & Worker Environment Tier
  - You can create multiple environments (dev, test, prod)

## Amazon S3

- It allows to store objects (files) in buckets (directories)
- Buckets must have a globally unique name.
- Buckets are defined at the region level.
- Naming convention
  - No uppercase
  - No underscore
  - 3-63 characters long
  - Not an IP
  - Must start with lowercase letter or number

## S3 Objects

- Objects (files) have a key - It is the full path of the file
- The key is composed of prefix + object name
  - s3://my-bucket/myfile.txt
  - s3://my-bucket/my-folder/another-folder/myfile.txt
- Bucket → my-bucket  
Prefix → my-folder/another-folder  
Object → myfile.txt
- There is no concept of 'directories' within buckets. The UI folder structure is for UX only  
Just keys with very long names that contain slashes (/)
- Object values are the content of the body
  - Max Object Size is 5TB (5000 GiB)
  - For uploading more than 5 GiB, use 'multi-part' upload
- Each object can have:
  - Metadata → list of key/value pairs - system or user metadata
  - Tags → Unicode key/value pairs - up to 10 useful for security/lifecycle
  - VersionID → If versioning is enabled

## S3 Versioning

- You can version your files on Amazon S3
- It is enabled at the bucket level
- Some key (file) overwrites will create increment version: 1, 2, 3...
- It is best practice to version your buckets
  - Protect against unintended deletes (ability to restore a version)
  - Easy roll back to previous version
- Any file that is not versioned prior to enabling versioning will have version 'null'
- Suspending/Disabling versioning does not delete the previous versions.

## S3 Encryption

- a) SSE-S3 : encrypts S3 objects using keys handled & managed by AWS
- b) SSE-KMS : leverage AWS Key Management Service to manage encryption keys
- c) SSE-C : when you want to manage your own encryption keys
- d) Client Side Encryption

### a) SSE-S3

- Object is encrypted server side
- AES-256 encryption type
- Must set header: 'x-amz-server-side-encryption': 'AES256'

### b) SSE-KMS

- Object is encrypted server side
- KMS advantages: user control & audit trail
- Must set header: 'x-amz-server-side-encryption': 'aws:kms'

### c) SSE-C

- Object is encrypted server side
- Amazon S3 does not store the encryption key you provide
- HTTPS must be used for uploading
- Encryption key must be provided on HTTP headers for every HTTP request made

## d) Client Side Encryption

- Client library such as the Amazon S3 Encryption client
- Clients must encrypt data themselves before sending to S3
- Clients must decrypt data themselves when retrieving from S3
- Customer fully manages the keys and encryption cycle

## S3 Security

- User based
  - IAM policies - which API calls should be allowed for a specific user from IAM console
- Resource based
  - Bucket Policies - bucket wide rules from the S3 console - allows cross account
  - Object Access Control List (ACL) - finer grain
  - Bucket Access Control List (ACL) - less common
- An IAM principal can access an S3 object if
  - The user IAM permissions allow it or the resource policy allows it
  - And there is no explicit Deny

## S3 Bucket Policies

### JSON based policy

- Resources : buckets and objects
- Actions : Set of API to Allow or Deny
- Effect : Allow / Deny
- Principal : Account or User to apply the policy to
- Use S3 bucket for policy to:
  - Grant public access to the bucket
  - Force objects to be encrypted at upload
  - Grant access to another account (Cross Account)

```
{  
  "Version": "2022-01-20",  
  "Statement": [  
    {  
      "Sid": "PublicRead",  
      "Effect": "Allow",  
      "Principal": "*",  
      "Action": [  
        "S3:GetObject"  
      ],  
      "Resource": [  
        "arn:aws:s3:::bucket/*"  
      ]  
    }  
  ]  
}
```

## Block Public Access

- Block public access to buckets and objects granted through
  - new access control lists (ACLs)
  - any access control lists (ACLs)
  - new public bucket or access point policies
- Block public and cross-account access to buckets and objects through any public buckets or access point policies.
  - These settings were created to prevent company data leaks
  - If you know your bucket should never be public, leave these on
  - It can be set at the account level.

## S3 - Security Options

### a) Networking

- Supports VPC endpoints (for instances in VPC without www internet)

### b) Logging and Audit

- S3 Access Logs can be stored in other S3 bucket
- API calls can be logged in AWS CloudTrail

### c) User Security

- MFA Delete : MFA (multi factor authentication) can be required in versioned buckets to delete objects
- Pre-Signed URLs : URLs that are valid only for a limited time
  - Eg: Premium video service for logged in users

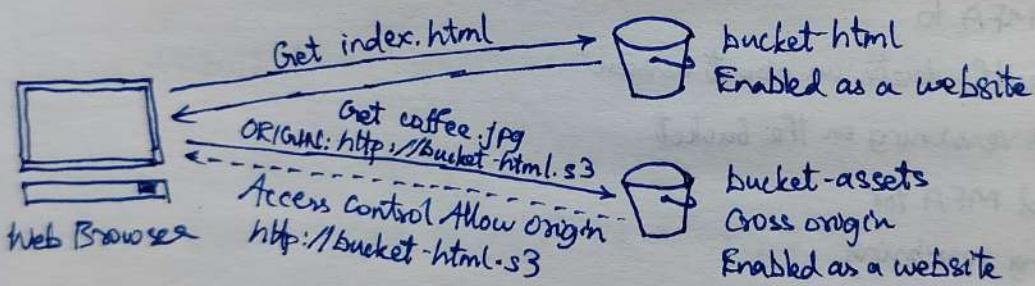
## S3 Websites

- S3 can host static websites and have them accessible on the www
- The website URL will be:
  - <bucket-name>.s3-website-<AWS-region>.amazonaws.com
  - <bucket-name>.s3-website.<AWS-region>.amazonaws.com
- If you get a 403 (Forbidden) error, make sure the bucket policy allows public reads

S3 → Properties → Static Website Hosting

## S3 CORS

- An origin is a scheme (protocol), host (domain) and port  
Eg: `https://www.example.com` (port is 443 for HTTPS)
- CORS → Cross-Origin Resource Sharing
- Web Browser based mechanism to allow requests to other origins while visiting the main origin.
  - Same origin: `https://example.com/app` & `https://example.com/app2`
  - Different origins: `https://www.example.com` & `https://other-example.com`
  - The requests won't be fulfilled unless the other origin allows for the requests using CORS Headers (ex: `Access-Control-Allow-Origin`)
  - If the client does a cross-origin request on our S3 bucket, we need to enable the correct CORS headers.
  - You can allow for a specific origin or for \* (all origins)



## AWS EC2 Instance Metadata

- It is powerful but one of the least known features to developers.
- It allows AWS EC2 instances to "know about themselves" without using an IAM Role for that purpose.
  - The URL is `http://169.254.169.254/latest/meta-data`
  - You can retrieve the IAM Role name from the metadata but you cannot retrieve the IAM policy
  - Metadata → Info about the EC2 instance
  - Userdata → Launch script of the EC2 instance

## AWS SDK

- To perform actions on AWS directly from your applications code (without using CLI)
- Official SDKs are: Java, .NET, Node.js, PHP, Python, Go, Ruby, Ett
- We have to use AWS Software Development kit when coding against AWS Services such as DynamoDB

such as DynamoDB

- The AWS CLI uses the Python SDK (boto 3)

• If you don't specify or configure a default region, then us-east-1 will be chosen by default

## S3 - MFA

- Multi-factor Authentication forces user to generate a code on a device (usually a mobile phone or hardware) before doing important operations on S3
- To use MFA-Delete, enable versioning on the S3 bucket
- You will need MFA to
  - Permanently delete an object version
  - Suspend versioning on the bucket
- You won't need MFA for
  - Enabling versioning
  - Listing deleted versions
- Only the bucket owner (root account) can enable/disable MFA-Delete
- MFA-Delete currently can only be enabled using the CLI

## S3 - Default Encryption

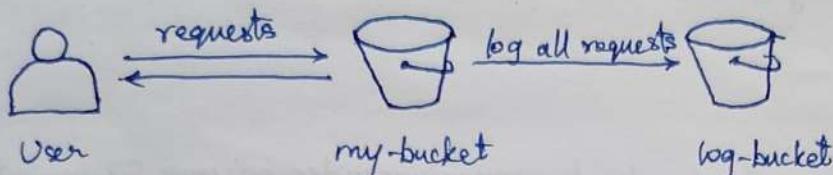
- One way to "force encryption" is to use a bucket policy and refuse any API call to PUT on S3 object without encryption headers
  - "S3:x-amz-server-side-encryption": "AES256"
- Another way is to use the "default encryption" option in S3. If you already uploaded an encrypted object it will not be re-encrypted
- Bucket Policies are evaluated before "default encryption"

S3 → Bucket → Properties → Default encryption

Automatically encrypt new objects stored in this bucket

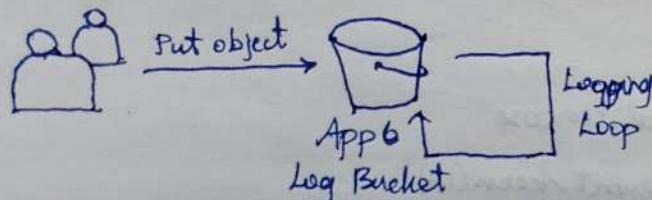
## S3 Access Logs

- For audit purpose, you may want to log all access to S3 buckets
- Any request made to S3, from any account, authorized or denied, will be logged into another S3 bucket
- The log data can be analyzed using data analysis tools or Amazon Athena



### Warning!

- Do not set your logging bucket to be monitored bucket as well. It will create a logging loop and your bucket will grow in size exponentially.

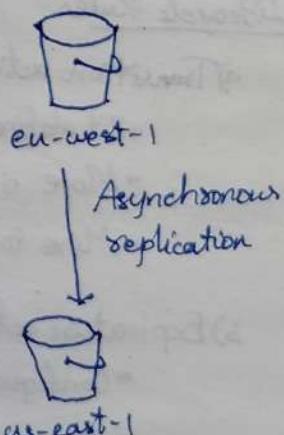


S3 → Appbucket → Properties → Server Access Logging

Note: By enabling S3 Access Logging, it will automatically update your bucket access control list (ACL) to include access to the S3 log delivery group.

## S3 Replication

- Must enable versioning in source and destination
- Cross Region Replication (CRR) & Same Region Replication (SRR)
- Buckets can be in different accounts. Copying is asynchronous
- Must give proper IAM permissions to S3
- CRR Use cases: compliance, lower latency access
- SRR Use cases: log aggregation, live replication
- After activating, only new objects are replicated (not retroactive)
- For DELETE operations:
  - Can replicate delete markers from source to target (optional setting)
  - Deletions with a version ID are not replicated (to avoid malicious deletes)
- There is no chaining of replication - if bucket 1 has replication into bucket 2, which has replication into bucket 3, then objects created in bucket 1 are not replicated to bucket 3



## S3 Pre-signed URL

- Can generate pre-signed URLs using SDK or CLI
- For downloads (easy - use CLI) and for uploads (hard - use SDK)
- Valid for a default of 3600 seconds, can change timeout with --expires-in [TimeSeconds]
- Users given a pre-signed URL inherit the permissions of the person who generated the URL for GET/PUT

## Examples

- Allow only logged in users to download a premium video on your S3 bucket.
- Allow an ever changing list of users to download files by generating URLs dynamically
- Allow temporarily a user to upload a file to a precise location in our bucket.

## S3 - Storage classes

- Amazon S3 Standard - General Purpose
- Amazon S3 Standard - Infrequent Access (IA)
- Amazon S3 One Zone - Infrequent Access
- Amazon S3 Intelligent Tiering
- Amazon Glacier
- Amazon Glacier - Deep Archive

## S3 Lifecycle Rules

### a) Transition actions

- It defines when objects are transitioned to another storage class
- Move objects to Standard IA class 60 days after creation
- Move to Glacier for archiving after 6 months

### b) Expiration actions

- Configure objects to expire (delete) after some time
- Access log files can be set to delete after a 365 days
- Can be used to delete old versions of files (if versioning is enabled)
- Can be used to delete incomplete multi-part uploads
- Rules can be created for a certain prefix (Eg: s3://mybucket/mp3/\*)
- Rules can be created for certain objects tags (Eg: Department: Finance)

## S3 Analytics

• You can setup S3 Analytics to help determine when to transition objects from Standard to Standard-IA

• Does not work for ONEZONE IA or GLACIER

• Report is updated daily - Takes about 24 h to 48 h for first start

• Good first step to put together Lifecycle Rules or to improve them

## S3 Baseline Performance

• Amazon S3 automatically scales to high request rates - latency 100-200 ms

• Your application can achieve atleast 3500 PUT/COPY/POST/DELETE and 5500 GET/HEAD requests per second per prefix in a bucket

• If you use SSE-KMS your S3 performance may be impacted by the KMS limits

• When you upload it calls the GenerateDataKey KMS API, when you download it calls the Decrypt KMS API

• Count towards the KMS quota per second (5500, 10000, 20000 req/s based on region)

• You can request a quota increase using Service Quotas Console

## S3 Enhanced Performance

### a) Multi-part upload

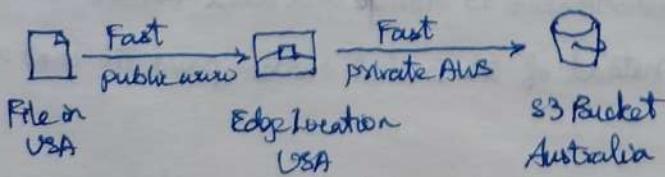
• Recommended for files > 100 MB, must use for files > 5 GB

• Can help parallelize uploads (speed up transfers)

### b) S3 Transfer Acceleration

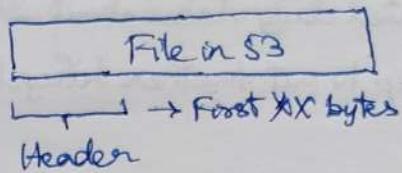
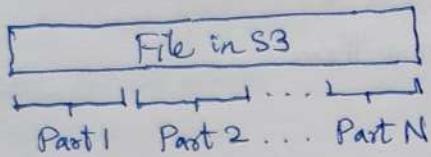
• Increase transfer speed by transferring file to an AWS edge location which will forward the data to the S3 bucket in the target region

• Compatible with multi-part upload



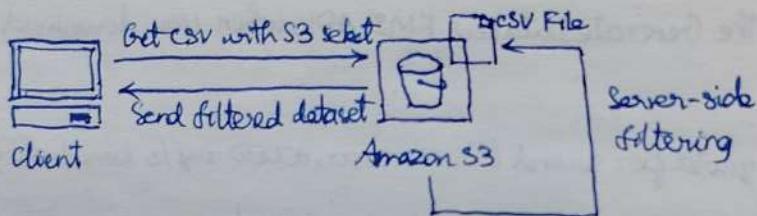
## S3 Byte-Range Fetches

- Parallelize GETs by requesting specific byte ranges.
- Better resilience in case of failures
- Can be used to speed up downloads (or) retrieve only partial data (head of a file)



## S3 Select & Glacier Select

- Retrieve less data using SQL by performing server-side filtering
- Can filter by rows & columns (simple SQL statements)
- Less network transfers less CPU cost client-side



## S3 Event Notifications

- S3:ObjectCreated, S3:ObjectRemoved, S3:ObjectRestore, S3:Replication
- Object name filtering is possible (\*.jpg)
- Can create as many S3 events as desired, typically delivers events in seconds but can sometimes take a minute or longer
  - If two writes are made to a single non-versioned object at the same time, it is possible that only a single event notification will be sent. To ensure for all writes, enable versioning.
- Targets are 1) SNS 2) SQS 3) Lambda

## Requester Pays

- In general, bucket owners pay for all Amazon S3 storage and data transfer costs associated
- With Requester Pays, the requester instead of the bucket owner pays the cost of the request and the data download from the bucket.
- Helpful when you want to share large datasets with other accounts.
- The requester must be authenticated on AWS (cannot be anonymous)

## Amazon Athena

- Serverless query service to perform analytics against S3 objects.
- Uses standard SQL language to query the files
- Supports CSV, JSON, ORC, Avro and Parquet (built on Presto)
- Pricing: \$5.00 per TB of scanned data
- Use compressed or columnar data for cost savings (less scan)
- Use cases: Business Intelligence/analytics reporting, analyze VPC Flow logs, ELB logs, CloudTrail, etc...

## AWS CloudFront

- Content Delivery Network (CDN)
- Improves read performance, content is cached at the edge
- 216+ Points of presence globally (edge locations)
- DDoS protection, integration with Shield, AWS Web Application Firewall
- Can expose external HTTPS and can talk to internal HTTPS backends

## CloudFront Geo Restriction

- Whitelist: Allow your users to access your content only if they're in one of the countries on a list of approved countries.
- Blacklist: Prevent your users from accessing your content if they're in one of the countries on a list of banned countries.
- The "country" is determined using a 3rd party Geo-IP database
- Use case: Copyright laws to control access to content

## CloudFront Signed URL / Signed Cookies

- You want to distribute paid shared content to premium users over the world
- We can use CloudFront Signed URL/Cookie. We attach a policy with:
  - Includes URL expiration
  - Includes IP ranges to access the data from
  - Trusted signers (which AWS accounts can create signed URLs)
- Signed URL → Access to individual files (one signed URL per file)
- Signed Cookies → Access to multiple files (one signed cookie for many files)

## CloudFront Signed URL vs S3 Pre-signed URL

### CloudFront Signed URL

- Allow access to a path, no matter the origin
- Account wide key-pair, only the root can manage it
- Can filter by IP, path, date expiration
- Can leverage caching features

### S3 Pre-Signed URL

- Issue a request as the person who pre-signed the URL
- Uses the IAM key of the signing IAM principal
- Limited lifetime

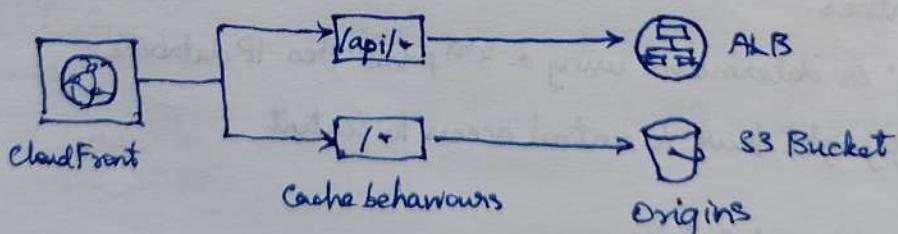
## CloudFront - Pricing

- CloudFront edge locations are all around the world. The cost of data out per edge location varies.
- You can reduce the no. of edge locations for cost reduction.
- Three Price classes:
  - i) Price class All : all regions - best performance
  - ii) Price class 200 : most regions but excludes the most expensive regions
  - iii) Price class 100 : only the least expensive regions

## CloudFront - Multiple Origin

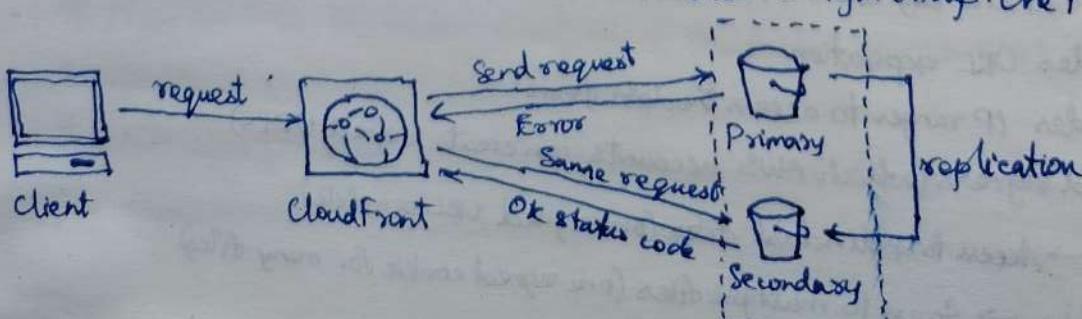
- To route different kind of origins based on the content type. Based on path pattern:

- /images/\*
- /api/\*
- /\*



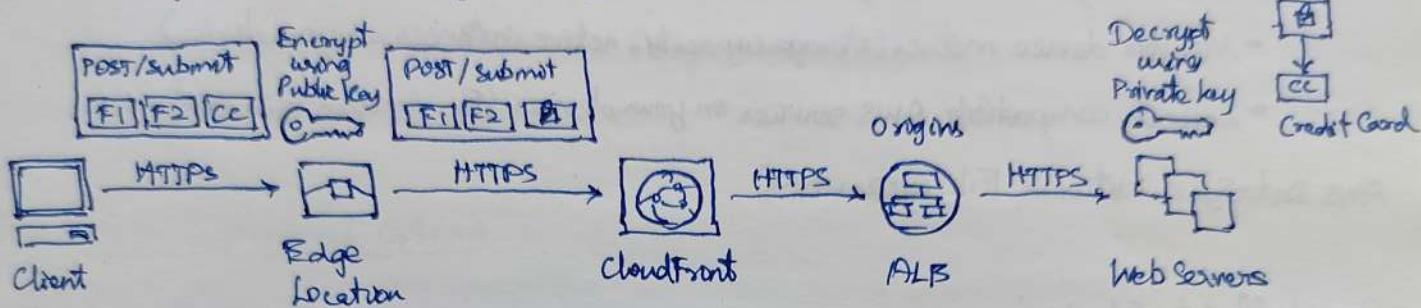
## CloudFront - Origin Groups

- To increase high availability and do failover. Origin Group: One Primary & One Secondary



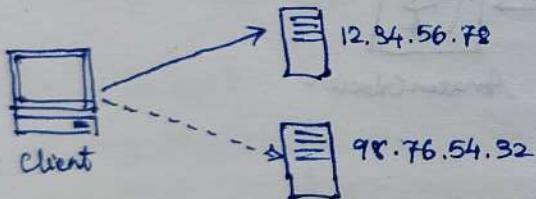
## CloudFront - Field Level Encryption

- Protect user sensitive information through application stack
- Adds an additional layer of security along with HTTPS
- Sensitive information encrypted at the edge closer to user. Uses asymmetric encryption.
- Usage:
  - Specify set of fields in POST requests that you want to be encrypted (upto 10 fields)
  - Specify the public key to encrypt them

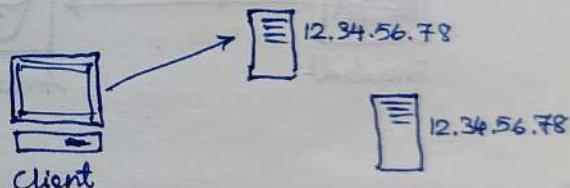


## Unicast IP vs Anycast IP

- One server holds one IP address
- Routing based on different address



- All servers hold same IP address
- Routed to the nearest one to client



## AWS Global Accelerator

- Leverages AWS internal network to route to your application. 2 anycast IP are created
- The Anycast IP send traffic directly to Edge locations and Edge locations send the traffic to App
- Works with Elastic IP, EC2 instances, ALB, NLB, public or private

### a) Consistent performance

- Intelligent routing to lowest latency and fast regional failover
- Internal AWS Network. No issue with client cache (because the IP doesn't change)

### b) Health checks

- Performs health check of applications - Create for disaster recovery
- Helps make your application global (failover less than 1 minute for unhealthy)

### c) Security

- Only 2 external IP need to be whitelisted. DDoS protection from AWS Shield

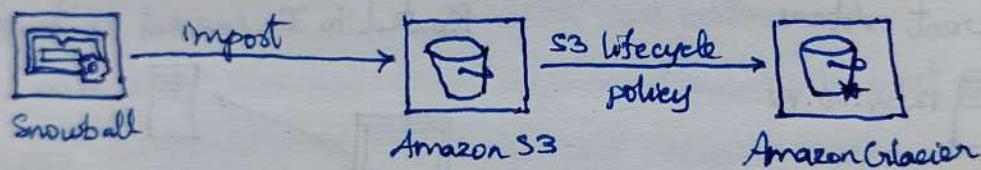
## AWS OPS HUB

- Historically to use Snow Family devices, you needed a CLI
- Now, you can use AWS Ops Hub - a software you install on your computer/laptop to manage your Snow Family Device
- Unlocking and configuring single or clustered devices
- Transferring files
- Launching and managing instances running on Snow Family devices
- Monitors device metrics (storage capacity, active instances on your device)
- Launch compatible AWS services on your devices (Eg: Amazon EC2 instances.)

AWS DataSync, Network File system (NFS)

## Snowball into Glacier

- Snowball cannot import to Glacier directly.
- You must use Amazon S3 first, in combination with an S3 Lifecycle policy.



## Amazon FSx

- Launch 3rd party high-performance file systems on AWS
- Fully managed service.
- a) FSx for Lustre
- b) FSx for Windows File Server
- c) FSx for NetApp ONTAP

## FSx for Lustre

- It is a type of parallel distributed file system for large scale computing, Linux + Cluster
- Machine learning, High performance computing, Video processing, Finance Modeling
- Scale up to 100s GB/s, millions of IOPS, sub-ms latencies
- Seamless integration with S3. Can read S3 as a file system and write back to S3
- Can be used from on-premise servers

## FSx for Windows

- EFS is a shared POSIX system for Linux systems.
- FSx for Windows is a fully managed Windows file system share drive
- Supports SMB protocol & Windows NTFS
- Microsoft Active Directory integration, ACLs, user quotas
- Built on SSD, scale up to 10s of GiB/s, millions of IOPS, 100s PB of data
- Can be accessed from your on-premises infrastructure
- Can be configured to be Multi-AZ (High availability)
- Data is backed-up daily to S3

## FSx File System Deployment Options

### a) Scratch File System

Temporary storage

- Data is not replicated (doesn't persist if file server fails)
- High burst (6x faster, 200 MBps per TiB)
- Usage: short-term processing, optimize costs

### b) Persistent File System

- Long-term storage
- Data is replicated within same AZ
- Replace failed files within minutes
- Usage: long-term processing, sensitive data

## Amazon FSx File Gateway

- Native access to Amazon FSx for Windows File Server
- local cache for frequently accessed data, lower latency
- Windows native compatibility (SMB, NTFS, Active Directory)
- Useful for group file shares and home directories

### File Gateway

- Configured S3 buckets are accessible using the NFS and SMB protocol
- Supports S3 standards, S3 IA, S3 One Zone IA
- Bucket access using IAM roles for each File Gateway
- Most recently used data is cached in the file gateway
- Can be mounted on many servers.
- Integrated with Active Directory (AD) for user authentication

### Volume Gateway

- Block storage using iSCSI protocol backed by S3
- Backed by EBS snapshots which can help restore on-premises volumes
- Cached volumes: low latency access to most recent data
- Stored volumes: entire dataset is on-premise, scheduled backups to S3

### Tape Gateway

- Some companies have backup processes using physical tapes
- With Tape Gateway, companies use the same processes but in the cloud
- Virtual Tape Library (VTL) backed by Amazon S3 and Glacier
- Backup data using existing tape-based processes (and iSCSI interface)
- Works with leading backup software vendors.

### Hardware Appliance

- Storage Gateway needs on-premises virtualization. You can use SG Hardware Appliance
- Works with File, Volume and Tape Gateway. It has the required CPU, memory, network, SSD cache resources. It is helpful for daily NFS backups in small data centers.

### Storage Gateway Summary

- File Gateway → File access / NFS - user auth with Active Directory (backed by S3)
- Volume Gateway → Volumes / Block storage / iSCSI (backed by S3 with EBS Snapshots)
- Tape Gateway → VTL Tape solution / Backup with iSCSI (backed up by S3 and Glacier)

## AWS Transfer Family

- A fully managed service for file transfers into and out of Amazon S3 or Amazon EFS using the FTP protocol
- Supported Protocols
  - AWS Transfer for FTP (File Transfer Protocol)
  - AWS Transfer for FTPS (File Transfer Protocol over SSL)
  - AWS Transfer for SFTP (Secure File Transfer Protocol)
- Managed infrastructure, Scalable, Reliable, Highly available (multi-AZ)
- Pay per provisioned endpoint per hour + data transfers in GB
- Store and manage user credentials within the service
- Integrate with existing authentication systems (Microsoft Active Directory, LDAP, Okta, Amazon Cognito, custom)
- Usage: sharing files, public datasets, CRM, ERP...

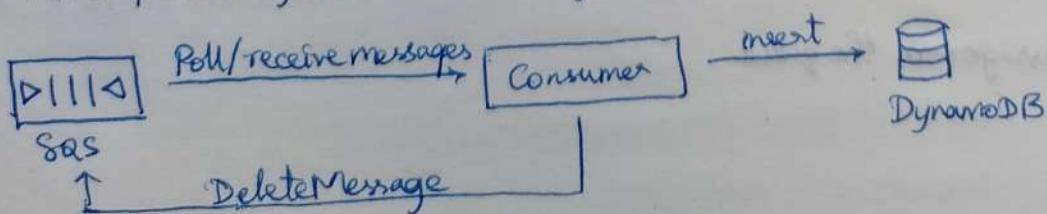


## SQS - Producing Messages

- Produced to SQS using the SDK (SendMessage API)
- The message is persisted in SQS until a consumer reads it and it gets deleted from queue.
- Message retention: default 4 days up to 14 days
- Eg: Send an order to be processed
  - Order id, customer id, any attribute you want

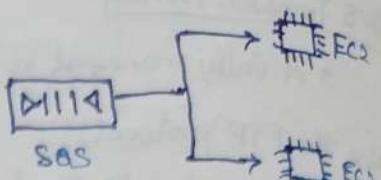
## SQS - Consuming Messages

- Consumers can be running on EC2 instance, servers or AWS Lambda
- Poll SQS for messages (receive upto 10 messages at a time)
- Process the messages (Eg. insert the message into DynamoDB)
- After processing, delete the messages using the DeleteMessage API

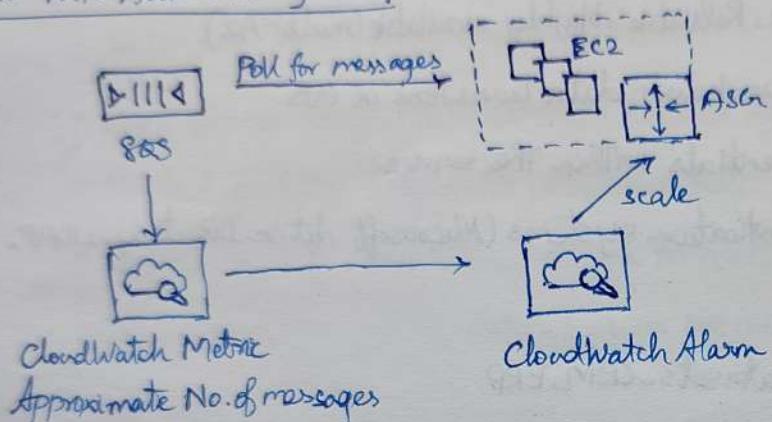


## SQS - Multiple EC2 Instances Consumers

- Consumers receive and process messages in parallel.
- At least once delivery
- Best-effort message ordering
- Consumers delete message after processing them
- We can scale consumers horizontally to improve throughput of processing



## SQS - With Auto Scaling Group



## Amazon SQS - Security (Same as SNS)

### a) Encryption

- In-flight encryption using HTTPS API
- At-rest encryption using KMS keys
- Client-side encryption if the client wants to do encryption/decryption himself

### b) Access Controls

- IAM policies to regulate access to the SQS API

### c) SQS Access Policies

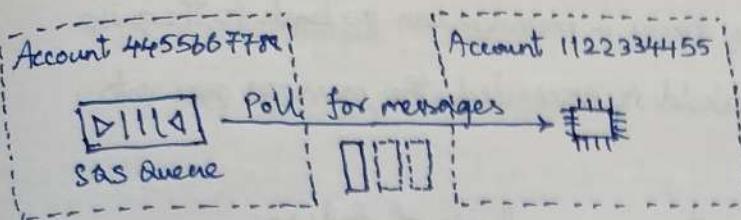
- Useful for cross-account access to SQS queues
- Useful for allowing other services (SNS, S3...) to write to an SQS queue

## Purge

- Delete all messages in the queue

## SQS Queue Access Policy

### Cross Account Access



"version": "2022-01-21"

"Statement": [

    "Effect": "Allow",

    "Principal": ["AWS": ["1122334455"]],

    "Action": ["sns:ReceiveMessage"],

    "Resource": "arn:aws:sns:us-east-1:",

        4455667788@queue"

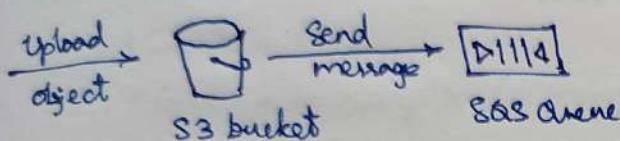
]

"Condition": [

    "ArnLike": {"aws:SourceArn": "arn:aws:s3:",

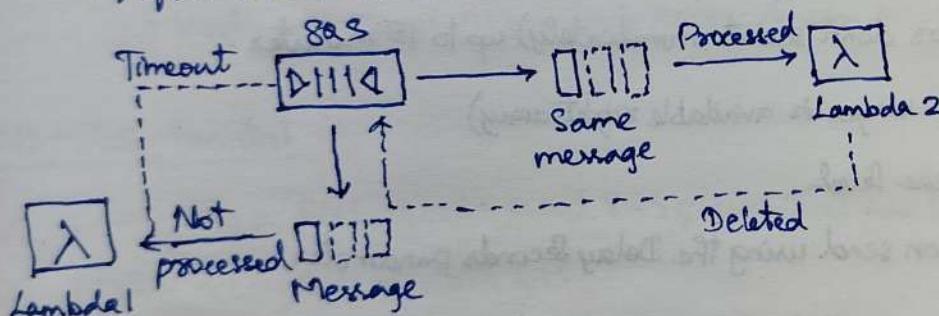
    "StringEquals": {"aws:SourceAccount": "n"}

### Publish S3 Event to SQS Queue



### SQS - Message Visibility Timeout

- After a message is polled by a consumer, it becomes invisible to other consumers.
- By default, the 'message visibility timeout' is 30 seconds, to be processed
- After the timeout is over, the message is 'visible' in SQS



- If a message is not processed within the visibility timeout it will be processed twice
- A consumer could call the `ChangeMessageVisibility` API to get more time
- If visibility timeout is high (hours) and consumers crashes, re-processing will take a lot of time

- If visibility timeout is too low (seconds), we may get duplicates

- So the visibility timeout should be set something reasonable for your application.

will take a lot of time

If visibility timeout is too low (seconds), we may get duplicates

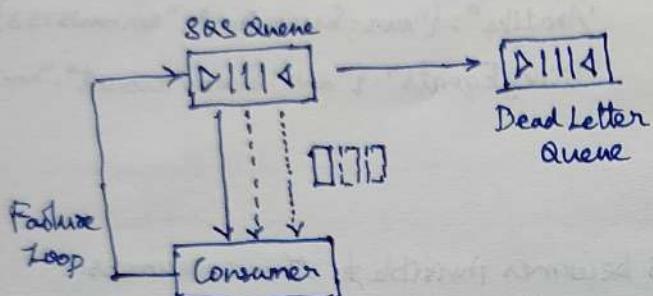
So the visibility timeout should be set something reasonable for your application.

## Amazon SQS - Dead Letter Queue

- If the consumer fails to process a message multiple times, it will create Failure Loop
- We can set a threshold of how many times a message can go back to the queue
- After the Maximum Receiver Threshold is exceeded, the message goes onto Dead Letter Queue (DLQ)

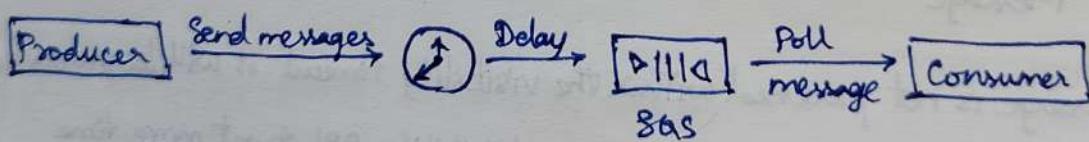
### Dead Letter Queue (DLQ)

- It is useful for debugging and root cause analysis of failure
- Make sure to process the messages in the DLQ before they expire:
- Good to set retention of 14 days in the DLQ



## SQS - Delay Queue

- Delay a message (consumers don't see it immediately) up to 15 minutes.
- Default is 0 seconds (messages are available right away)
- Can set a default at queue level
- Can override the default on send using the Delay Seconds parameter



## SQS - Long Polling

- When a consumer requests messages from the queue, it can optionally wait for messages to arrive if there are none in the queue
- Long Polling decreases the no. of API calls made to SQS while increasing the efficiency and latency of your application.
- The wait time can be between 1 sec to 20 sec (20 sec preferable)
- Long Polling is preferable to Short Polling
- Long Polling can be enabled at the queue level or at the API level using WaitTimeSeconds

## SQS - FIFO queue

- FIFO → First In First Out (Ordering of messages in the queue)
- Limited throughput: 300 msg/s without batching, 3000 msg/s with batching
- Exactly-once send capability (by removing duplicates)
- Messages are processed in order by the consumer

## AWS SNS - How to publish

### a) Topic Publish (using the SDK)

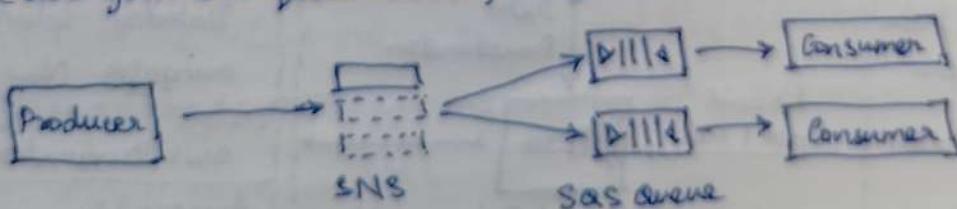
- Create a topic
- Create a subscription (or many)
- Publish to the topic

### b) Direct Publish (for mobile apps SDK)

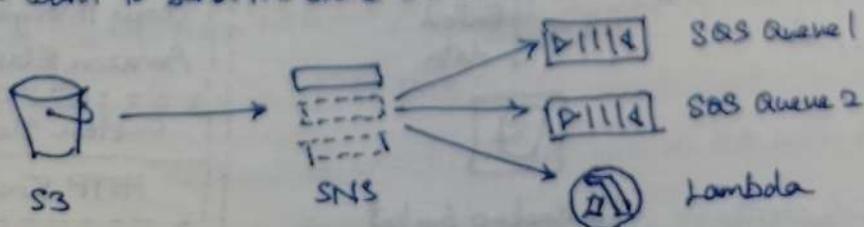
- Create a platform application
- Create a platform endpoint
- Publish to the platform endpoint
- Works with Google GCM, Apple APNS, Amazon ADM...

## SNS + SQS : Fan Out

- Push once in SNS, receive in all SQS queues that are subscribers
- Fully decoupled, no data loss
- SQS allows for: data persistence, delayed processing and retries of work
- Ability to add more SQS subscribers over time
- Make sure your SQS queue access policy allows for SNS to write

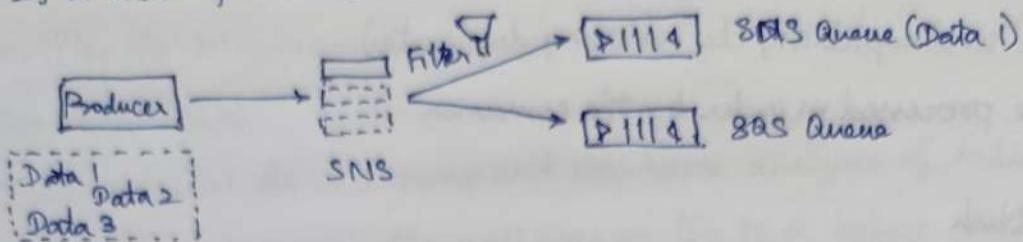


- If you want to send the same S3 event to many SQS queues, we fan-out

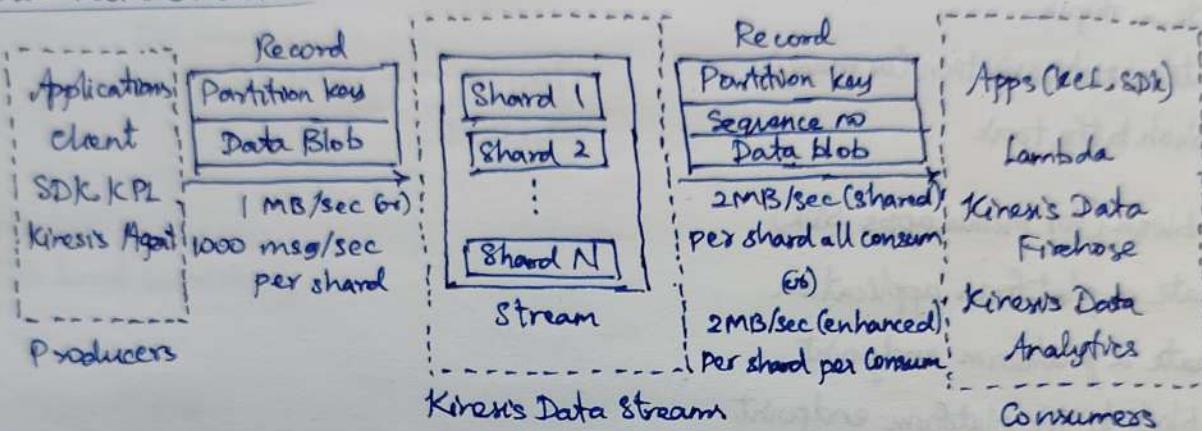


## SNS - Message Filtering

- JSON policy used to filter messages sent to SNS topic's subscriptions
- If a subscription doesn't have a filter policy, it receives every message

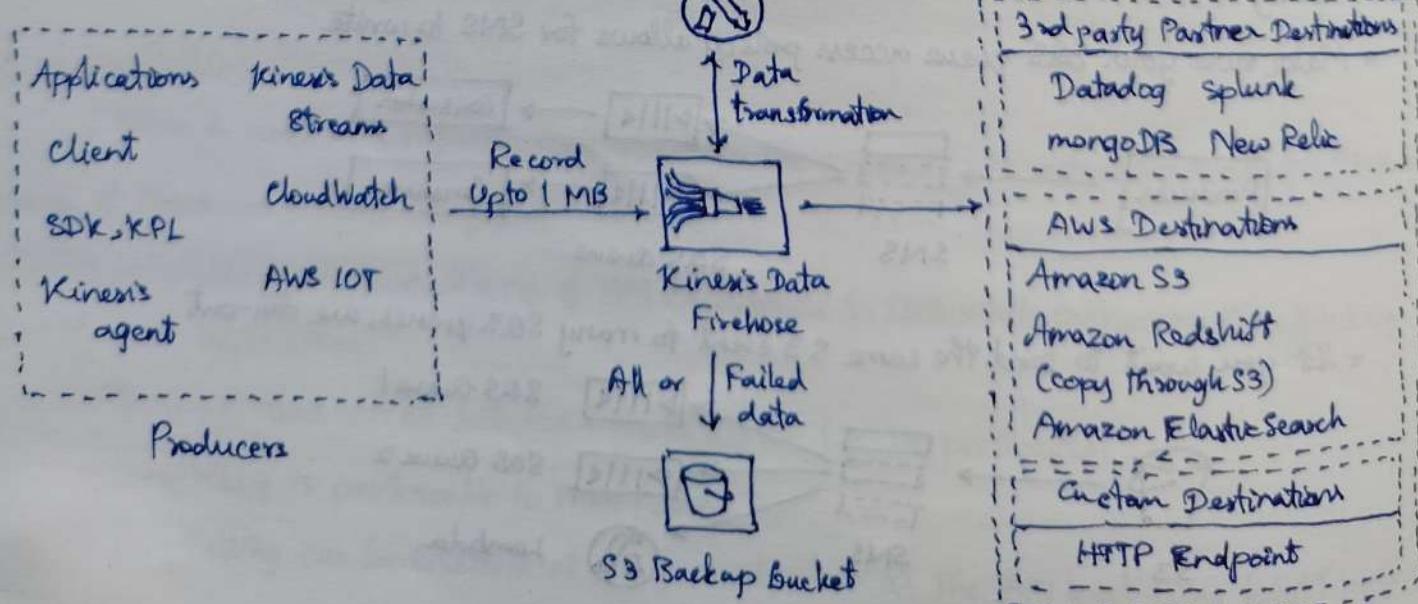


## Kinesis Data Streams



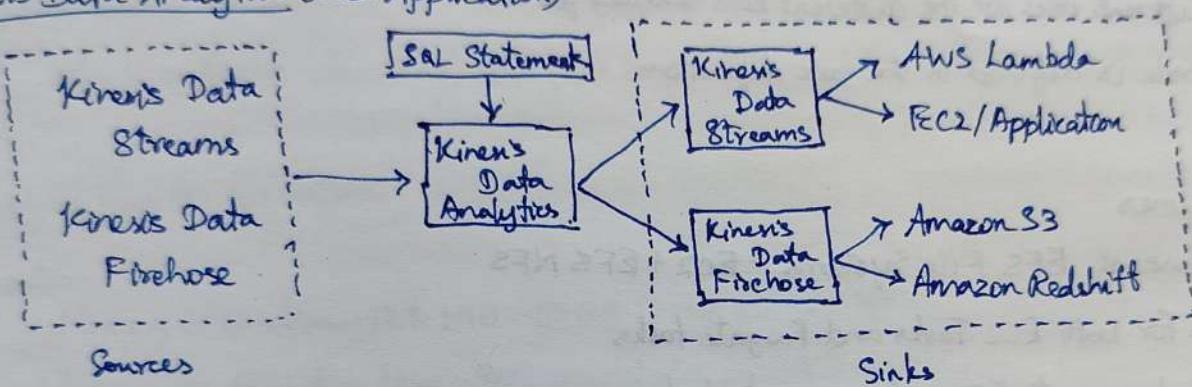
- Billing is per shard provisioned, can have as many shards as you want
- Retention between 1 day (default) to 365 days
- Ability to reprocess (replay) data
- Once data is inserted in kinesis, it can't be deleted (immutability)
- Data that shares the same partition goes to the same shard (ordering)

## Kinesis Data Firehose



- Kinesis Data Firehose is a fully managed service, no administration, automatic scaling
- Pay for data going through Firehose
- Near Real Time
  - 60 seconds latency minimum for non full batches
  - Or minimum 32 MB of data at a time
- Supports many data formats, conversions, transformations, compression
- Supports custom data transformations using AWS Lambda
- Can send failed or all data to a backup S3 Bucket

### Kinesis Data Analytics (SQL Application)



- Perform real time analytics on kinesis streams using SQL
- Fully managed, no servers to provision. Automatic scaling.
- Real-time analytics, pay for actual consumption rate
- Can create streams out of the real-time queries
- Use cases: Time series analytics, Real-time dashboards, Real-time metrics

### Amazon MQ

- SQS & SNS are 'cloud-native' services and they are using proprietary protocols from AWS
- Traditional applications running from on-premise may use open protocols such as MQTT, AMQP, STOMP, Openwire, WS
- When migrating from on-prem to cloud instead of re-engineering the application to use SQS & SNS, we can use Amazon MQ → Managed Apache ActiveMQ
  - It doesn't scale as much as SQS/SNS
  - It runs on a dedicated machine, can run in HA with failover
  - It has both queue feature (~SQS) and topic features (~SNS)

## IAM Roles for ECS Tasks

### a) EC2 Instance Profile

- Used by the ECS agent
- Makes API calls to ECS service
- Send container logs to CloudWatch Logs
- Pull Docker image from ECR
- Reference sensitive data in Secrets Manager or SSM Parameter Store

### b) ECS Task Role

- Allow each task to have a specific role
- Use different roles for the different ECS services you run
- Task role is defined in the task definition

## ECS Data Volumes

- Make use of EFS File Systems - EC2 + EFS NFS
- Works for both EC2 Tasks and Fargate tasks
- Ability to mount EFS volumes onto tasks
- Tasks launched in any AZ will be able to share the same data on the EFS volume.
- Fargate + EFS = serverless + data storage without managing servers.
- Use Case: persistent multi-AZ shared storage for your containers

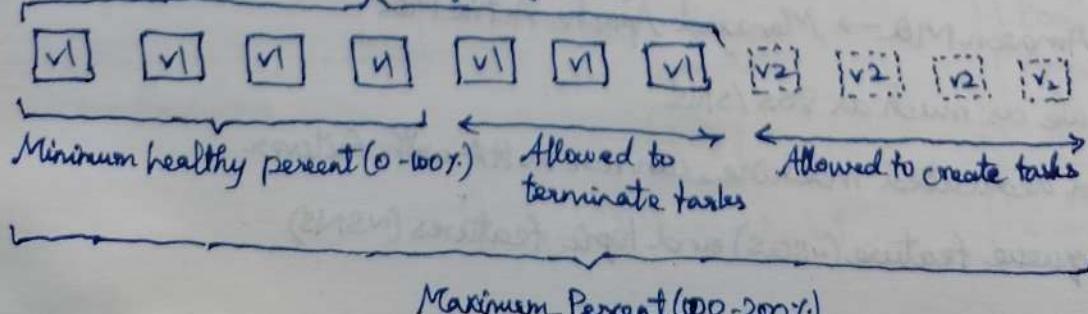
## ECS Rolling Updates

- When updating from v1 to v2, we can control how many tasks can be started and stopped and in which order.
- When you have an ECS update, you will have 2 settings:

i) Minimum healthy percent = 100

ii) Maximum percent = 200

Actual Running capacity (v1, v2)



## Amazon EKS

- Amazon Elastic Kubernetes Service
- It's a way to launch managed Kubernetes clusters on AWS
- Kubernetes is an open-source system for automatic deployment, scaling and management of containerized (usually Docker) application
  - It's an alternative to ECS, similar goal but different API
  - EKS supports EC2 if you want to deploy worker nodes or Fargate to deploy serverless containers
    - Use Case: If your company is already using Kubernetes on-premises or in another cloud, and wants to migrate to AWS using Kubernetes
    - Kubernetes is cloud-agnostic (can be used in any cloud - Azure, GCP)

## AWS Lambda Limits - per region

### a) Execution

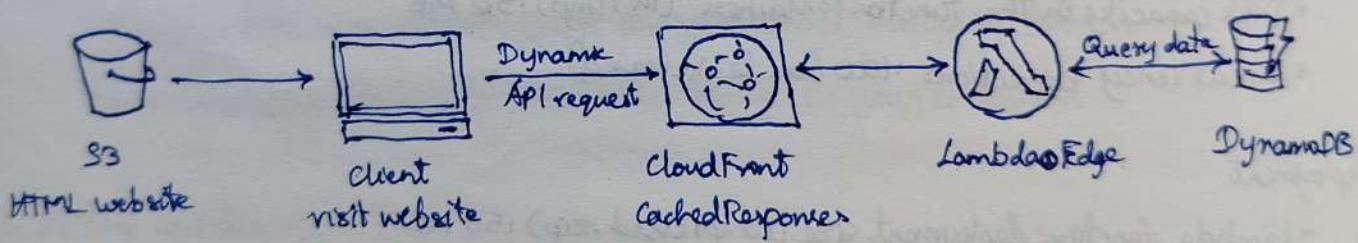
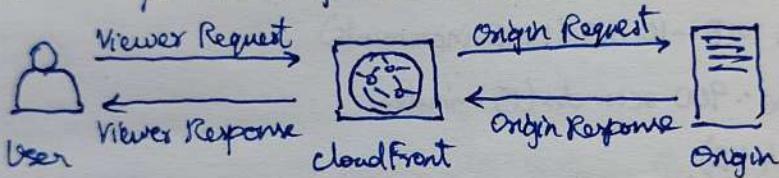
- Memory allocation: 128 MB - 10 GiB (1 MB increments)
- Maximum execution time: 900 seconds (15 minutes)
- Environment variables (4 KB)
- Disk capacity in the 'function container' (in /tmp): 512 MB
- Concurrency executions: 1000 (can be increased)

### b) Deployment

- Lambda function deployment size (compressed .zip): 50 MB
- Size of uncompressed deployment (node dependencies): 250 MB
- Can use the /tmp directory to load other files at startup
- Size of environment variables: 4 KB

## Lambda@Edge

- You have deployed a CDN using CloudFront. What if you wanted to run a global AWS Lambda alongside? Or how to implement request filtering before reaching your application?
- Use Lambda@Edge to deploy Lambda functions alongside your CloudFront CDN
  - Build more responsive applications
  - You don't manage servers. Lambda is deployed globally
  - Customize the CDN content
- You can change CloudFront requests and responses
  - i) Viewer request → After CloudFront receives a request from a viewer
  - ii) Origin request → Before CloudFront forwards the request to the origin
  - iii) Origin response → After CloudFront receives the response from the origin
  - iv) Viewer response → Before CloudFront forwards the response to the viewer
- You can also generate responses to viewers without ever sending the request to the origin



## Use Cases:

- Website Security and Privacy
- Dynamic web applications at the Edge
- Search Engine Optimization (SEO)
- Intelligently Route across origins and data centers
- Bot mitigation at the Edge
- Real time Image Transformation
- A/B Testing
- User Authentication and Authorization
- User Prioritization
- User tracking and Analytics

## DynamoDB

- The database is already created. We have to just create tables.
- Each table has a Primary Key (must be decided at creation time)
- Each table can have infinite number of items (rows)
- Each item has attributes (can be added over time - can be null)
- Maximum size of an item is 400 kB
- Data types supported are:
  - Scalar Types - String, Number, Binary, Boolean, Null
  - Document Types - List, Map
  - Set Types - String Set, Number Set, Binary Set

## DynamoDB - Read/Write Capacity Modes

- Control how you manage your table's capacity (read/write throughput)

### a) Provisioned Mode (default)

- You specify the number of reads/writes per second
- You need to plan capacity beforehand
- Pay for provisioned Read Capacity Units (RCU) & Write Capacity Units (WCUs)
- Possibility to add auto-scaling mode for RCU & WCU

### b) On-Demand Mode

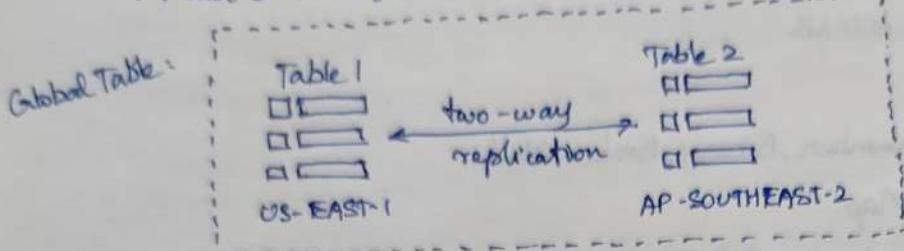
- Read/writes automatically scale up/down with your workloads
- No capacity planning needed
- Pay for what you use, more expensive (\$\$\$)
- Great for unpredictable workloads

## DynamoDB Streams

- Ordered stream of item-level modifications (create/update/delete) in a table.
- Stream records can be sent to Kinesis Data Streams, read by AWS Lambda & KCL App
- Data retention for up to 24 hours
- Use Cases: react to changes in real-time (welcome email to users)
  - Analytics
  - Insert into derivative tables or into Elasticsearch
  - Implement cross-region replication

## DynamoDB Global Table

- Make a DynamoDB table accessible with low latency in multiple regions
- Active-Active replication
- Applications can Read and Write to the table in any region
- Must enable DynamoDB Streams as a pre-requisite



## DynamoDB Time-to-Live (TTL)

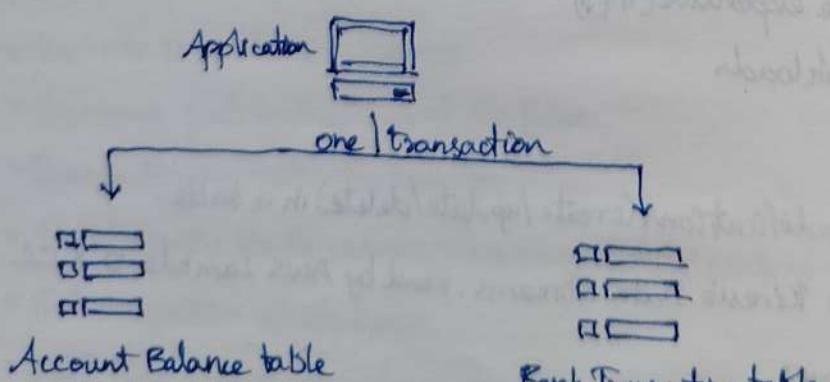
- Automatically delete items after an expiry timestamp
- Use cases: reduce stored data by keeping only current items, adhere to regulatory obligations

## DynamoDB Indexes

- Global Secondary Index (GSI) or Local Secondary Index (LSI)
- Allow to query on attributes other than the Primary key
- With Indexes, we can query by other attributes of the items

## DynamoDB Transactions

- Transaction is written to both tables or none



### Read Modes:

- Eventual Consistency
- Strong Consistency
- Transactional

### Write Modes:

- Standard
- Transactional

- Provides Atomicity, Consistency, Isolation and Durability (ACID)
- Consumes 2x WCVs & RCVs
- Use cases: financial transactions, managing orders

## API Gateway - Endpoint Types

- a) Edge-Optimized (default) : For global clients
  - Requests are routed through the CloudFront Edge Locations (improves latency)
  - The API Gateway still lives in only one region
- b) Regional
  - For clients within the same region
  - Could manually combine with CloudFront (more control over the caching strategies and the distribution)
- c) Private
  - Can only be accessed from your VPC using an interface VPC endpoint (ENI)
  - Use a resource policy to define access

## API Gateway - Security

### IAM Permissions - For users within AWS

- Create an IAM policy authorization and attach to User / Role
- API Gateway verifies IAM permissions passed by the calling application
- Good to approve access within your own infrastructure
- Leverages 'Sig v4' capabilities where IAM credential are in headers

### Lambda Authorizer (Custom Authorizer)

- Uses AWS Lambda to validate the token in header being passed. OAuth/SAML/3rd party
- Option to cache result of authentication.
- Lambda must return an IAM policy for the user

### Cognito User Pools

- Cognito fully manages user lifecycle
- API Gateway verifies identity automatically from AWS Cognito
- No custom implementation required
- Cognito only helps with Authentication, not Authorization
- Authentication - Cognito User Pools  
Authorization - API Gateway Methods

## AWS Cognito

We want to give our users an identity so that they can interact with our application.

### a) Cognito User Pools:

- Signin functionality for app users
- Integrate with API Gateway

### b) Cognito Identity Pools (Federated Identity)

- Provide AWS Credentials to users so they can access AWS resources directly
- Integrate with Cognito User Pools as an identity provider

### c) Cognito Sync

- Synchronize data from device to Cognito. Store preferences, configuration, state of app
- Might be deprecated and replaced by AppSync. Offline capability, requires FIP not CIP

## AWS Cognito User Pool (CUP)

- Create a serverless database of user for your mobile apps
- Simple login: Username (or email) / password combination
- Possibility to verify emails/phone numbers and add MFA
- Can enable Federated Identities (Facebook, Google, SAML).
- Sends back a JSON Web Token (JWT)
- Can be integrated with API Gateway for authentication and Application Load Balancer

## Federated Identity Pools

- Provide direct access to AWS Resources from Client side
- Log in to federated identity provider or remain anonymous
- Get temporary AWS credentials back from the Federated Identity Pool
- These credentials come with a pre-defined IAM policy stating their permissions
- Eg: Provide (temporary) access to write to S3 bucket using Facebook login

## AWS SAM

- Serverless Application Model
- Framework for developing and deploying serverless applications
- All the configuration is YAML code
  - Lambda Functions
  - DynamoDB tables
  - API Gateway
  - Cognito User Pools
- SAM can help you to run Lambda, API Gateway, DynamoDB locally
- SAM can use CodeDeploy to deploy Lambda functions

## Redshift

- Data is loaded from S3, DynamoDB, DMS, other DBs
- From 1 node to 128 nodes, up to 128 TB of space per node
- Leader node: for query planning, results aggregation
- Compute node: for performing the queries, send results to leader
- Redshift Spectrum: perform queries directly into S3 (no need to load)
- Backup & Restore, Security, VPC / IAM / KMS, Monitoring
- Redshift Enhanced VPC Routing: Copy/unload goes through VPC

## Snapshots & DR

- Redshift has no 'Multi-AZ' mode
- Snapshots are point-in-time backups of a cluster, stored internally in S3
- Snapshots are incremental (only what has changed is saved)
- You can restore a snapshot into a new cluster
- Automated: every 8 hours, every 5 GB, or on a schedule. Set retention
- Manual: snapshot is retained until you delete it.
- You can configure Amazon Redshift to automatically copy snapshots (automated or manual) of a cluster to another AWS Region

## ElasticSearch / OpenSearch

- You can search any field, even partially matches
- It's common to use ElasticSearch as a complement to another database
- It also has some usage for Big Data applications
- You can provision a cluster of instances
- Built-in integrations: Amazon Kinesis Data Firehose, AWS IoT and Amazon CloudWatch Logs
- Security through Cognito & IAM, KMS encryption, SSL & VPC
- Comes with Kibana (visualization) & Logstash (log ingestion) - ELK stack
- Eg: In DynamoDB, you can find only by primary key or indexes. With ElasticSearch you can search any field

→ ElasticSearch = Search / Indexing

## CloudWatch Custom Metrics

- Possibility to define and send your own custom metrics to CloudWatch
- Eg: memory (RAM) usage, disk space, number of logged in users
- Use API call PutMetricData
- Ability to use dimensions (attributes) to segment metrics
  - Instance.id
  - Environment.name
- Metric resolution (StorageResolution API parameter - two possible values)
  - Standard: 1 minute (60 seconds)
  - High Resolution: 1/5/10/30 second(s) - Highest cost
- Important: Accepts metric data points two weeks in the past and two hours in the future  
(make sure to configure your EC2 instance time correctly)

## CloudWatch Dashboards

- Custom dashboards for quick access to key metrics and alarms. Dashboards are global.
- It can include graphs from different AWS accounts and regions. You can change the time range of the dashboards. You can setup automatic refresh (10s, 1m, 2m, 5m, 15m)
- It can be shared with people who don't have an AWS account.  
(public, email address, 3rd party SSO provider through Amazon Cognito)

## CloudWatch Logs

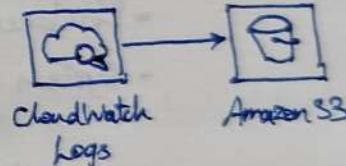
- Log groups : arbitrary name usually representing an application
- Log stream : instances within application / log file / containers
- Can define log expiration policies (never expire, 30 days, etc...)
- It can send logs to:
  - Amazon S3 (exports)
  - AWS Lambda
  - Kinesis Data Streams
  - ElasticSearch
  - Kinesis Data Firehose

## CloudWatch - Metric Filter & Insights

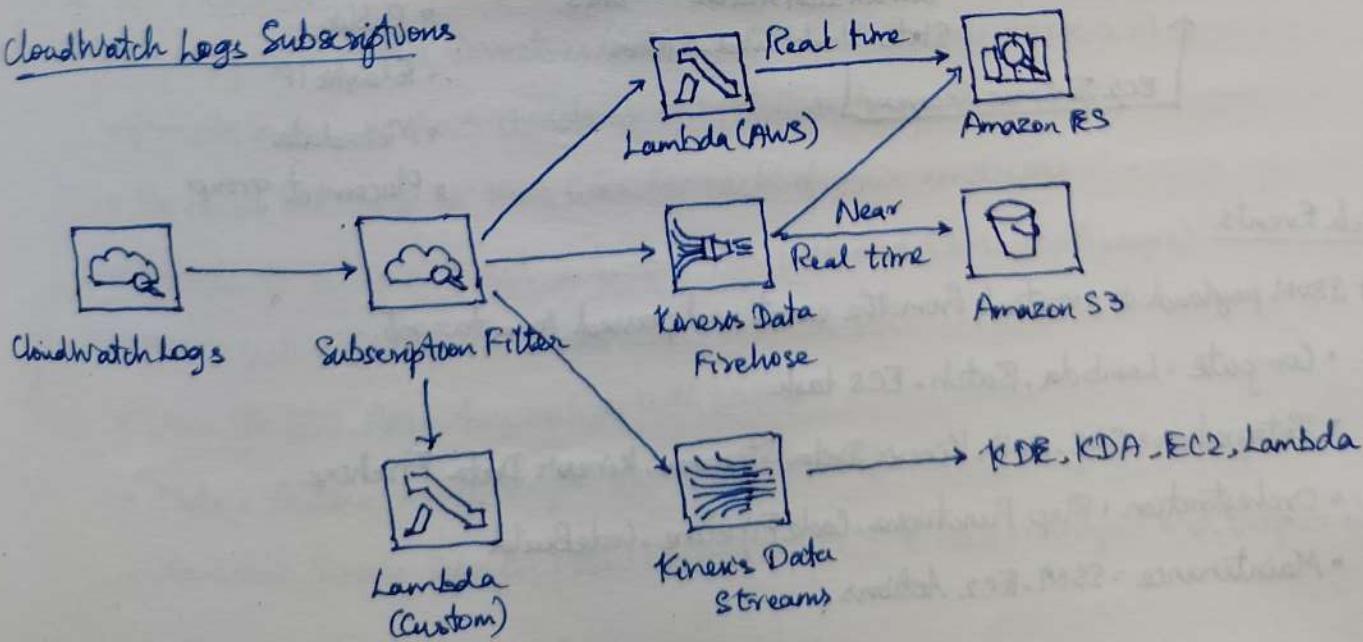
- CloudWatch Logs can use filter expressions
  - For eg: find a specific IP inside of a log
  - Or count occurrences of 'Error' in your logs
- Metric filters can be used to trigger CloudWatch alarms
- Logs Insights can be used to query logs and add queries to CloudWatch Dashboards

## CloudWatch-S3 Export

- Log data can take up to 12 hours to become available for export
- The API call is Create Export Task
- Not near-real time or real time, use logs Subscriptions instead

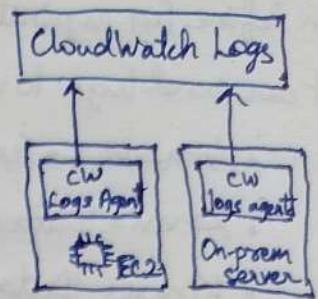


## CloudWatch Logs Subscriptions



## CloudWatch Logs for EC2

- By default, no logs from your EC2 machine will go to CloudWatch.
- You need to run a CloudWatch agent on EC2 to push the log files you want.
- Make sure IAM permissions are correct
- CloudWatch log agent can be setup on-premises too



### a) CloudWatch Logs Agent

- Old version of the agent
- Can only send to CloudWatch Logs

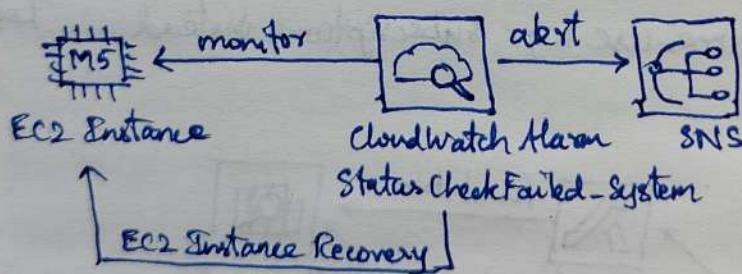
### b) CloudWatch Unified Agent

- Collect additional system level metrics such as RAM, processes, etc...
- Collect logs to send to CloudWatch Logs
- Centralized configurations using SSM Parameter Store

## EC2 Instance Recovery

### • Status Check:

- Instance status → check the EC2 VM
- System status → check the underlying hardware



### Recovery

- Save Private
- Public
- Elastic IP
- Metadata
- Placement group

## CloudWatch Events

- A JSON payload is created from the event and passed to a target.
  - Compute: Lambda, Batch, ECS task
  - Integration: SQS, SNS, Kinesis Data Streams, Kinesis Data Firehose
  - Orchestration: Step Functions, CodePipeline, CodeBuild
  - Maintenance: SSM, EC2 Actions

## Amazon EventBridge Schema Registry

- EventBridge can analyze the events on your bus and infer the schema
- The schema registry allows you to generate code for your application that will know in advance how data is structured in the event bus.
- Schema can be versioned.
- EventBridge allows extension to add event buses for your custom applications and your third-party SaaS apps.

## Identity Federation in AWS

- Federation lets users outside of AWS to assume temporary role for accessing AWS resources.
- These users assume identity provided access role.
- Federation can have many flavors:
  - SAML 2.0
  - Custom Identity Broker
  - Web Identity Federation w/wo Amazon Cognito
  - Single Sign-on
  - Non SAML with AWS Microsoft AD
- Using federation, you don't need to create IAM users (user management is outside of AWS)

## SAML 2.0 Federation

- To integrate Active Directory / ADFS with AWS (or any SAML 2.0)
- Provide access to AWS Console or CLI (through temporary creds)
- No need to create an IAM user for each of your employees
- Needs to setup a trust between AWS IAM and SAML (both ways)
- SAML 2.0 enables web-based cross domain SSO
- Use the STS API: AssumeRoleWithSAML
- Note: federation through SAML is the 'old way' of doing things
- Amazon Single Sign-on (SSO) Federation is the new managed and simpler way.

## Custom Identity Broker Application

- Use only if identity provider is not compatible with SAML 2.0
- The identity broker must determine the appropriate IAM policy
- Uses the STS API: AssumeRole or GetFederationToken

## IAM Conditions

- aws:SourceIP: restrict the client IP from which the API calls are being made
  - "aws:SourceIp": [  
    "192.0.2.0/24",  
    "203.0.113.0/24"]
- aws:RequestedRegion: restrict the region the API calls are made to
  - "aws:RequestedRegion": [  
    "eu-central-1",  
    "eu-east-1"]
- Restrict based on tags
  - "ec2:ResourceTag/Project": "~~",
  - "aws:PrincipalTag/Department": "~~"
- Force MFA
  - "BoolIfExists": {"aws:MultiFactorAuthPresent": false}

## IAM Roles vs Resource Based Policies

- When you assume a role (user, application or service) you give up your original permissions and take the permissions assigned to the role.
- When using a resource based policy, the principal doesn't have to give up his permissions
- Eg: User on Account A needs to scan a DynamoDB table in Account A and dump it in an S3 bucket on Account B
- Supported by: Amazon S3 buckets, SNS topics, SQS queues.

## IAM Permission Boundary

- IAM Permission Boundaries are supported by users and roles (not groups)
- Advanced feature to use a managed policy to set the maximum permissions an IAM entity can get.

(Example)

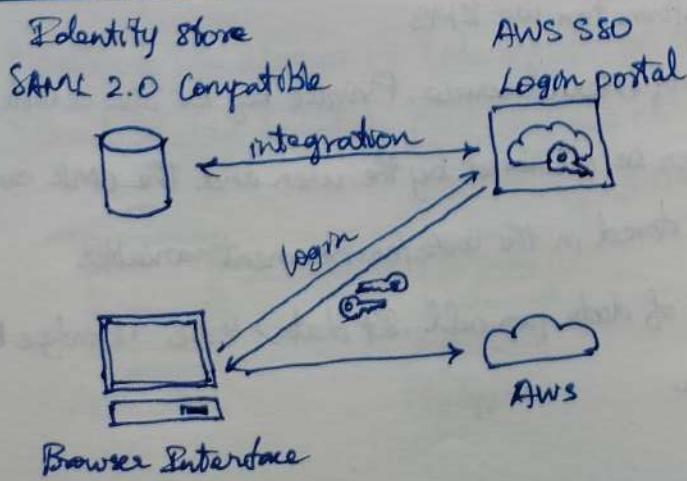
"Action": [ "s3:\*", "cloudwatch\*", "ec2:\*" ] + "Action": "iam:CreateUser" = [No permission]

## IAM Permission Boundary

## AWS Resource Access Manager (RAM)

- Share AWS resources that you own with other AWS accounts
- Share with any account or within your organization
- Avoid resource duplication
- VPC Subnets
  - Allow to have all the resources launched in the same subnets
  - Must be from the same AWS Organizations
  - Cannot share security groups and default VPC
  - Participants can share their own resources in there
  - Participants cannot view, modify, delete resources that belong to others or owner
- AWS Transit Gateway, Route53 Resolver Rules, License Manager Configurations

## AWS Single-Sign-On (SSO)



## KMS - Customer Master Key (CMK) Types

### a) Symmetric (AES-256 keys)

- First offering of KMS - single encryption key that is used to Encrypt and Decrypt
- AWS services that are integrated with KMS use Symmetric CMKs
- Necessary for envelope encryption
- You never get access to the key unencrypted (must call KMS API to use)

### b) Asymmetric (RSA & ECC keypairs)

- Public (Encrypt) and Private key (Decrypt) pair
- Used for Encrypt/Decrypt or Sign/Verify operations
- The public key is downloadable but you can't access the Private key unencrypted
- Use Case: encryption outside of AWS by users who can't call the KMS API

## AWS KMS

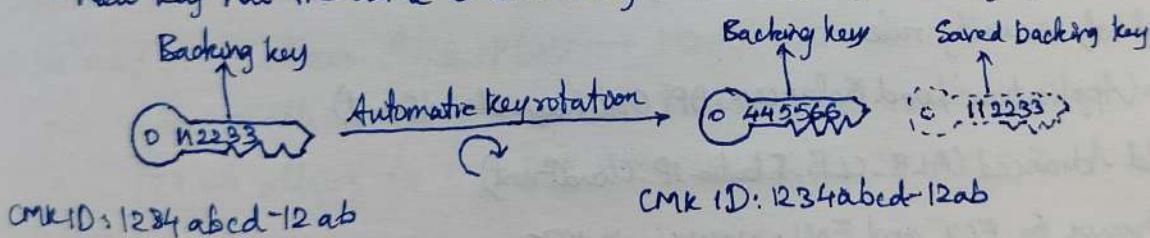
- Able to fully manage the keys & policies
  - Create
  - Rotation policies
  - Disable / Enable
- Able to audit key usage (using CloudTrail)
- Three types of Customer Master Keys (CMK):
  - AWS Managed Service Default CMK - free
  - User Keys created in KMS : \$1/month
  - User Keys imported (must be 256-bit symmetric key) : \$1/month
- + pay for API call to KMS (\$0.03/10000 calls)
- Anytime you need to share sensitive information use KMS
  - Database passwords, Credentials of external service, Private key for SSL certificates
- The CMK used to encrypt data can never be retrieved by the user and the CMK can be rotated for extra security. Encrypted keys can be stored in the code/environment variables
- KMS can help only in encrypting 4 KB of data per call. If data > 4 KB → Envelope Encryption
- KMS keys are bound to a specific region

## KMS Key Policies

- Control access to KMS keys. 'similar' to S3 bucket policies
- Difference: you cannot control access without them
- Default KMS Key Policy:
  - Created if you don't provide a specific KMS Key Policy
  - Complete access to the IAM policies to the KMS key
  - Complete access to the key to the root user → entire AWS account.
- Custom KMS Key Policy:
  - Define users, roles that can access the KMS key
  - Define who can administer the key
  - Useful for cross-account access of your KMS key

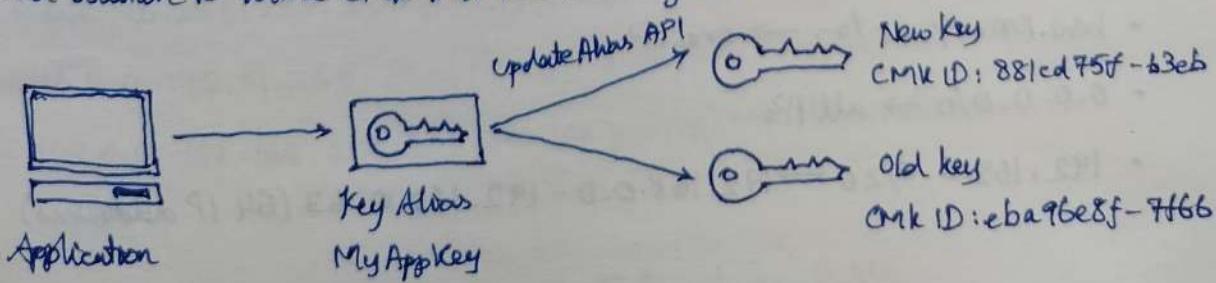
## KMS Key Rotation

- For Customer-managed CMK (not AWS managed CMK)
- If enabled: automatic key rotation happens every 1 year
- Previous key is kept active so you can decrypt old data
- New key has the same CMK ID (only the backing key is changed)



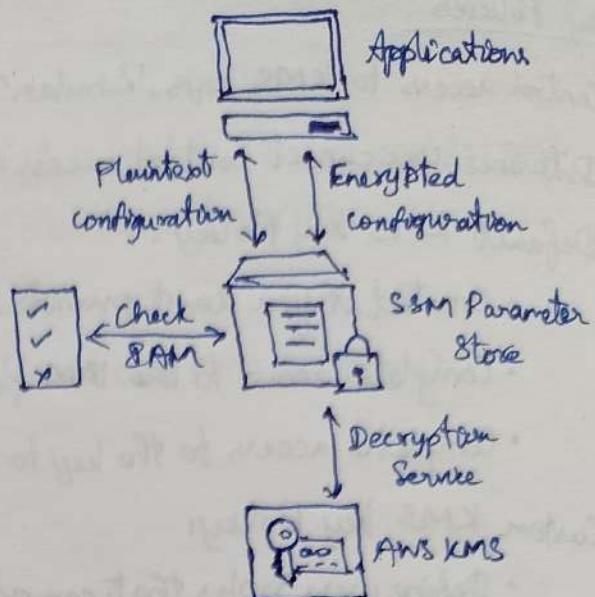
## KMS Manual Key Rotation

- When you want to rotate key every 90 days, 180 days, etc...
- New key has a different CMK ID. Keep the previous key active so you can decrypt old data
- Better to use aliases (to hide the change of key for the application)
- Good solution to rotate CMK that are not eligible for automatic rotation (asymmetric CMK)



## SSM Parameter Store

- Secure storage for configuration and secrets
- Optional Seamless Encryption using KMS
- Serverless, Scalable, durable, easy SDK
- Version tracking of configuration /secrets
- Configuration management using Path & IAM
- Notifications with CloudWatch events
- Integration with CloudFormation



## Parameter Policies

- It is for advanced parameters, not for standard parameters
- Allow to assign a TTL to a parameter (expiration date) to force updating or deleting sensitive data such as passwords
- Can assign multiple policies at a time

## AWS Firewall Manager

- Manage rules in all accounts of an AWS organization
- Common set of security rules
- WAF rules (Application Load Balancer, API Gateway, CloudFront)
- AWS Shield Advanced (ALB, CLB, Elastic IP, CloudFront)
- Security Groups for EC2 and ENI resources in VPC

## CIDR - IPv4

- Classless Inter-Domain Routing - a method for allocating IP addresses
- Used in Security Groups rules and AWS networking in general
- They help to define an IP address range.
  - $xx.xx.yy.zz/32 \rightarrow$  one IP
  - $0.0.0.0/0 \rightarrow$  all IPs
  - $192.168.0.0/26 \rightarrow 192.168.0.0 - 192.168.0.63$  (64 IP addresses)

## Understanding CIDR - IPv4

- A CIDR consists of two components

- Base IP
  - Represents an IP contained in the range (xx.xxx.xxx.xx)
  - Eg: 10.0.0.0, 192.168.0.0

- Subnet Mask
  - Defines how many bits can change in the IP
  - Eg: /0 - /32
  - Can take two forms:

/8 → 255.0.0.0

/16 → 255.255.0.0

/24 → 255.255.255.0

/32 → 255.255.255.255

Octets
IP: $\underline{1^{\text{st}}} \cdot \underline{2^{\text{nd}}} \cdot \underline{3^{\text{rd}}} \cdot \underline{4^{\text{th}}}$
/32 → no octet can change
/24 → last octet can change
/16 → last 2 octets can change
/8 → last 3 octets can change
/0 → all octets can change

## Subnet Mask

- Basically allows part of the underlying IP to get additional next values from the base IP

192.168.0.0/32 → allows for 1 IP ( $2^0$ ) → 192.168.0.0

192.168.0.0/31 → allows for 2 IP ( $2^1$ ) → 192.168.0.0 - 192.168.0.1

• 192.168.0.0/30 → allows for 4 IP ( $2^2$ ) → 192.168.0.0 - 192.168.0.3

• 192.168.0.0/29 → allows for 8 IP ( $2^3$ ) → 192.168.0.0 - 192.168.0.7

## Public vs Private IP (IPv4)

- The Internet Assigned Numbers Authority (IANA) established certain blocks of IP addresses for the use of private (LAN) and public (Internet) addresses.

Private IP can only allow certain values:

• 10.0.0.0 - 10.255.255.255 (10.0.0.0/8) → In Big Networks

• 172.16.0.0 - 172.31.255.255 (172.16.0.0/12) → AWS default VPC

• 192.168.0.0 - 192.168.255.255 (192.168.0.0/16) → Home networks

## Public IP

- All the rest of the IP addresses on the Internet are Public

## Default VPC

- All new AWS accounts have a default VPC
- New EC2 instances are launched onto the default VPC if no subnet is specified
- Default VPC has Internet connectivity and all EC2 instances inside it have public IPv4 addresses
- We also get a public and private IPv4 DNS names.

## VPC in AWS - IPv4

- VPC → Virtual Private Cloud
- You can have multiple VPCs in an AWS region (max 5 per region - soft limit)
- Max CIDR per VPC is 5, for each CIDR:
  - Min size is /28 (16 IP addresses)
  - Max size is /16 (65536 IP addresses)
- Because VPC is private, only the Private IPv4 ranges are allowed
- Your VPC CIDR should not overlap with your other networks (Eg: corporate)

## VPC - Subnet IPv4

- AWS reserves 5 IP addresses (first 4 & last 1) in each subnet
- These 5 IP addresses are not available for use and can't be assigned to an EC2 instance
- Example: If CIDR block 10.0.0.0/24, then reserved IP addresses are
  - 10.0.0.0 → Network Address
  - 10.0.0.1 → reserved by AWS for the VPC router
  - 10.0.0.2 → reserved by AWS for mapping to Amazon provided DNS
  - 10.0.0.3 → reserved by AWS for future use
  - 10.0.0.255 → Network Broadcast Address. AWS does not support broadcast in a VPC  
Therefore the address is reserved.

• Note: If you need 29 IP addresses for EC2 instances:

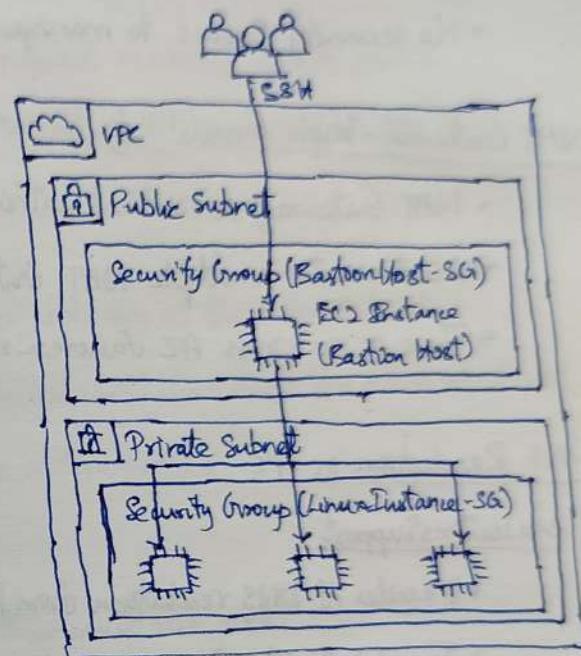
- X You can't choose a subnet of size /27 (32 IP addresses,  $32 - 5 = 27$ )
- ✓ You need to choose a subnet of size /26 (64 IP addresses,  $64 - 5 = 59$ )

## Internet Gateway (IGW)

- Allows resources (e.g. EC2 instances) in a VPC connect to the Internet
- It scales horizontally and is highly available and redundant
- Must be created separately from a VPC
- One VPC can only be attached to one IGW and vice versa
- Internet Gateways on their own do not allow Internet access. Route table must also be edited.

## Bastion Hosts

- We can use a Bastion Host to SSH into our private EC2 instances.
- The bastion is in the public subnet which is then connected to all other private subnets.
- Bastion host security group must be tightened.
- Note: Make sure the bastion host only has port 22 traffic from the IP address you need, not from the security groups of your other EC2 instances



## NAT Instance

- NAT → Network Address Translation
- Allows EC2 instances in private subnets to connect to the Internet
- Must be launched on a public subnet
- Must disable EC2 setting: Source/Destination Check
- Must have Elastic IP attached to it
- Route tables must be configured to route traffic from private subnets to the NAT instance

## Comments

- Pre-configured Amazon Linux AMI is available - Reached end of support on Dec 31, 2020
- Not highly available/resilient group setup out of the box
- Internet traffic bandwidth depends on EC2 instance type
- You must manage Security Groups & rules

## NAT Gateway

- AWS-managed NAT, higher bandwidth, high availability, no administration
- Pay per hour for usage and bandwidth
- NATGW is created in a specific Availability zone, uses an Elastic IP
- Can't be used by EC2 instance in the same subnet (only from other subnets)
- Requires an IGW (Private Subnet → NATGW → IGW)
- 5 Gbps of bandwidth with automatic scaling up to 45 Gbps
- No security Groups to manage/required.

## NAT Gateway - High Availability

- NAT Gateway is resilient within a single Availability Zone
- Must create multiple NAT Gateways in multiple AZs for fault-tolerance
- There is no cross AZ failover needed because if an AZ goes down it doesn't need NAT

## DNS Resolution in VPC

### enableDnsSupport

- Decides if DNS resolution from Route 53 Resolver server is supported for the VPC
- True (default) - it queries the Amazon Provider DNS Server at 169.254.169.253 or the reserved IP address at the base of the VPC IPv4 network range plus two (-2)

### enableDnsHostnames

- By default,
  - True → default VPC
  - False → newly created VPCs
- Won't do anything unless enableDnsSupport = true
- If True, assigns public hostname to EC2 instance if it has a public IPv4



EC2 Instance

Private DNS : ip-private-ipv4-address.ec2.internal

Public DNS : ec2-public-ipv4-address.compute-1.amazonaws.com

## NACL

- Network Access Control Lists are like a firewall which control traffic from and to subnets
- One NACL per subnet, new subnets are assigned the Default NACL
- NACL Rules:
  - Rules have a number (1-32766), higher precedence with a lower number
  - First rule match will drive the decision
  - Eg: If you define #100 Allow 10.0.0.10/32 & #200 Deny 10.0.0.10/32  
The IP address will be allowed because 100 has a higher precedence over 200
  - The last rule is an asterisk (\*) and denies a request in case of no rule match
  - AWS recommends adding rules by increment of 100, just so that new rules can be added inbetween.
- Newly created NACLs will deny everything
- NACL are a great way of blocking a specific IP address at the subnet level
- Default NACL accepts everything inbound/outbound with the subnets it's associated with
- Do not modify the default NACL, instead create custom NACLs

## Ephemeral Ports

- For any two endpoints to establish a connection, they must use ports.
- Client connects to a defined port, and expects a response on an ephemeral port
- Different Operating Systems use different port ranges, examples:
  - ZANA & MS Windows 10 → 49152-65535
  - Many Linux Kernels → 32768-60999

## VPC - Reachability Analyzer

- A network diagnostics tool that troubleshoots network connectivity between two endpoints in VPC
- It builds a model of the network configuration, then checks the reachability based on these configurations (it doesn't send packets)
- When the destination is
  - Reachable - it produces hop-by-hop details of the virtual network path
  - Not reachable - it identifies the blocking component(s). Eg: Configuration issues.
- Use cases: Troubleshoot connectivity issues, ensure network configuration is as intended

## VPC Flow Logs Syntax

2 12345678 eni-1235b8 172.31.6.139 172.31.16.21 20641 22620 4249 ACCEPT ok  
 ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓  
 version AccountID Interface-ID srcaddr destaddr srcport destport protocol packets bytes logby

- srcaddr & destaddr → help identify problematic IP
- srcport & destport → help identify problematic ports
- Action - success or failure of the request due to Security Group/NACL
- Can be used for analytics on usage patterns or malicious behaviors
- Query VPC flow logs using Athena on S3 or CloudWatch Logs insights

## AWS Site-to-Site VPN

### a) Virtual Private Gateway (VGW)

- VPN concentrator on the AWS side of the VPN connection
- VGW is created and attached to the VPC from which you want to create the Site-to-Site VPN connection
- Possibility to customize the ASN (Autonomous System Number)

### b) Customer Gateway (CGW)

- Software applications or physical device on customer side of the VPN connection

## Customer Gateway Device (on-premises)

What IP address to use?

- Public Internet-routable IP address for your Customer Gateway Device

If it's behind a NAT device that's enabled for NAT traversal (INAT-T) use the public IP address of the NAT device

Important: Enable Route Propagation for the virtual Private Gateway in the route table that is associated with your subnets

Note: If you need to ping your EC2 instances from on-premises, make sure you add the ICMP protocol on the inbound of your security groups

## AWS VPN CloudHub

- Provide secure connection/communication between multiple sites, if you have multiple VPN connections
- Low cost hub and spoke model for primary or secondary network connectivity between different locations (VPN only)
- It's a VPN connection so it goes over the public Internet
- To set it up - connect multiple VPN connections on the same VGW setup dynamic routing and configure route tables.

## Direct Connect (DX)

- Provides a dedicated private connection from a remote network to your VPC.
- Dedicated connection must be setup between your DC and AWS Direct Connect locations.
- You need to setup a Virtual Private Gateway on your VPC
- Access public resources (S3) and private (EC2) on same connection
- Supports both IPv4 and IPv6
- Use cases:
  - Increase bandwidth throughput - working with large data sets - lower cost
  - More consistent network experience - applications using real-time data feeds
  - Hybrid Environments (on-prem + cloud)
- If you want to setup a Direct Connect to one or more VPC in many different regions (same account), you must use a Direct Connect Gateway

## Connection Types

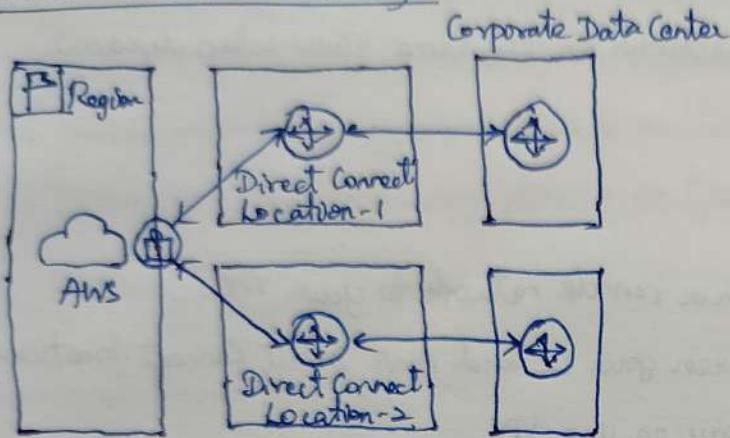
- Dedicated Connections: 1 Gbps and 10 Gbps capacity
  - Physical ethernet port dedicated to a customer
  - Request made to AWS first, then completed by AWS Direct Connect Partners.
- Hosted Connections: 500 Mbps, 50 Mbps, to 10 Gbps
  - Connection requests are made via AWS Direct Connect Partners
  - Capacity can be added or removed on demand
  - 1, 2, 5, 10 Gbps available at select AWS Direct Connect Partners.

Note: For both the connections the lead times are often longer than 1 month to establish a new connection

## Direct Connect - Encryption

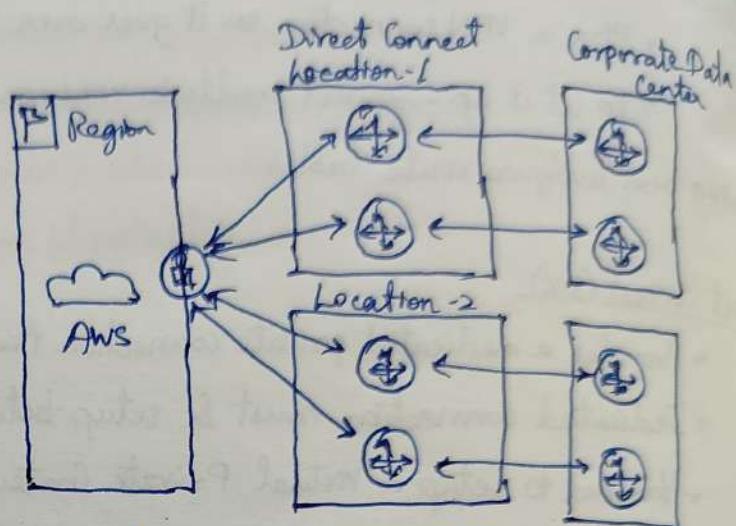
- Data in transit is not encrypted but is private
- AWS Direct Connect + VPN provides an IPsec-encrypted private connection.
- Good for an extra level of security but slightly more complex to put in place

## Direct Connect - Resiliency



One connection at multiple locations

HIGH RESILIENCY FOR CRITICAL WORKLOADS

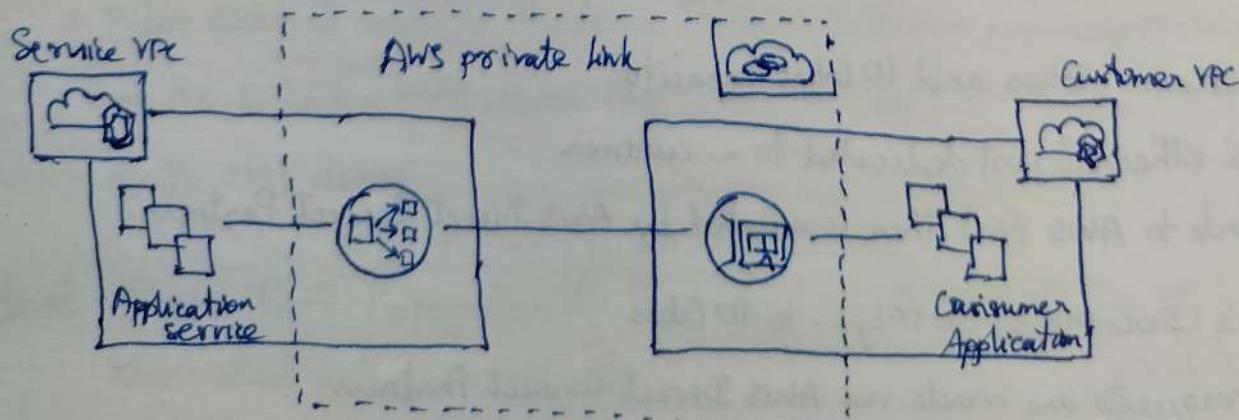


Separate connections in more than one location

MAXIMUM RESILIENCY FOR CRITICAL WORKLOADS

## AWS PrivateLink (VPC Endpoint Services)

- Most secure and scalable way to expose a service to 1000s of VPCs (own or other accounts)
- Does not require VPC peering, Internet gateway, NAT, route tables
- Requires a network load balancer (Service VPC) and ENI (Customer VPC) or GWLB
- If the NLB is in multiple AZ, and the ENIs in multiple AZ, the solution is fault tolerant



## EC2-Classic & AWS ClassicLink (deprecated)

- EC2-Classic → instances run in a single network shared with other customers
- ClassicLink → Allows you to link EC2-Classic instances to a VPC in your account.
  - Must associate a security group
  - Enables communication using private IPv4 addresses
  - Removes the need to make use of public IPv4 addresses or Elastic IP addresses.

## Transit Gateway: Site to Site VPN ECMP

- ECMP → Equal cost multi path routing
- Routing strategy to allow to forward a packet over multiple best path
- Use case: create multiple Site-to-Site VPN connections to increase the bandwidth of your connection to AWS

## VPC-Traffic Mirroring

- Allow you to capture and inspect network traffic in your VPC
- Route the traffic to security appliances that you manage
- Capture the traffic
  - From (Source) - ENIs
  - To (Targets) - an ENI or a Network Load Balancer
- Capture all packets or capture the packets of your interest (optionally truncate packets)
- Source and Target can be in the same VPC or different VPCs (VPC Peering)
- Use cases: content inspection, threat monitoring, troubleshooting...

## IPv6

- IPv6 is successor of IPv4. It is designed to provide  $3.4 \times 10^{38}$  unique IP addresses
- Every IPv6 address is public and Internet-routable (no private range)
- Format → x.x.x.x.x.x.x.x (x is hexadecimal, range can be from 0000 to ffff)
- Examples:
  - 2001:db8:3333:4444:5555:6666:7777:8888
  - 2001:db8:3333:4444:cccc:dddd:eeee:ffff
  - :: → all 8 segments are zero
  - 2001:db8:: → last 6 segments are zero | ::1234:5678 → first 6 segments are zero
  - 2001:db8::1234:5678 → middle 4 segments are zero

## IPv6 in VPC

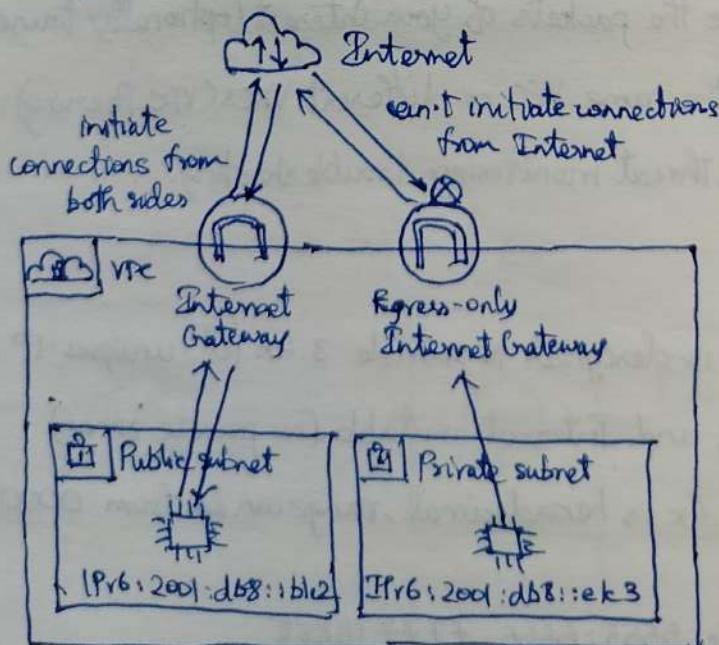
- IPv4 cannot be disabled for your VPC and subnets
- You can enable IPv6 (they're public IP addresses) to operate in dual-stack mode.
- Your EC2 instances will get at least a private internal IPv4 and a public IPv6
- They can communicate using either IPv4 or IPv6 to the Internet through an Internet Gateway

## IPv6 Troubleshooting

- If you cannot launch an EC2 instance in your subnet
  - It's not because it cannot acquire an IPv6 (the space is very large)
  - It's because there are no available IPv4 in your subnet
- Solution: create a new IPv4 CIDR in your subnet

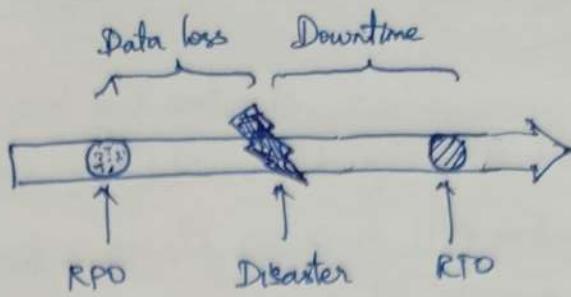
## Egress-only Internet Gateway

- Similar to NAT Gateway but for IPv6
- Allows instances in your VPC outbound connections over IPv6 while preventing the internet to initiate an IPv6 connection to your instances.
- You must update the Route Tables.



## Disaster Recovery Strategies

- a) Backup and Restore
- b) Pilot Light
- c) Warm standby
- d) Hot site/Multi-site



RPO:  
Recovery Point Objective

RTO:  
Recovery Time Objective

### b) Pilot Light

- A small version of the app is always running in the cloud
- Use for critical core.
- Faster than Backup and Restore as critical systems are already up

### c) Warm Standby

- Full system is up and running but at minimum size
- Upon disaster, we can scale to production load.

### d) Multi Site/Hot Site

- Very low RTO (minutes or seconds) - very expensive
- Full production scale is running AWS and on-premise

## DMS - Database Migration Service

- Quickly and securely migrate databases to AWS, resilient, self-healing
- The source database remains available during the migration
- Supports:
  - Homogeneous migrations. ex. Oracle to Oracle
  - Heterogeneous migrations. ex. Microsoft SQL Server to Aurora
- Continuous Data replication using CDC
- You must create an EC2 instance to perform the replication tasks.

## AWS Schema Conversion Tool (SCT)

- Convert your Database's Schema from one engine to another
- Eg: OLTP : (SQL Server or Oracle) to MySQL, PostgreSQL, Aurora
- Eg: OLAP : (Teradata or Oracle) to Amazon Redshift
- You do not need to use SCT if you are migrating the same DB engine

## On-Premise strategy with AWS

- Ability to download Amazon Linux 2 AMI as a VM (-isoformat)
  - VMware, KVM, VirtualBox(Oracle VM), Microsoft Hyper-V
- VM Import/Export
  - Migrate existing applications into EC2
  - Create a DR repository strategy for your on-premise VMs
  - Can export back the VMs from EC2 to on-premise
- AWS Application Discovery Service
  - Gather information about your on-premise servers to plan a migration
  - Server utilization and dependency mappings
  - Track with AWS Migration Hub
- AWS Database Migration Service (DMS)
  - Replicate on-premise → AWS, AWS → AWS, AWS → on-premise
- AWS Server Migration Service (SMS)
  - Incremental replication of on-premise live servers to AWS

## AWS DataSync

- Move large amount of data from on-premise to AWS
- Can synchronize to: Amazon S3 (any storage classes - including Glacier), EFS, FSx
- Move data from your NAS or file system via NFS or SMB
- Replication tasks can be scheduled hourly, daily, weekly
- Leverage the DataSync agent to connect to your systems
- Can setup a bandwidth limit